

Project on CPU, GGPU, and TPU

Lixian Chen, Zhanhao Zhang, Jingbin Cao

4/14/2021

Preparing Required Packages

```
data <- read.csv("../data/Runtime.csv")
head(data)
```


```
##      Runtime Processor MatrixSize MatrixOperation Trial
## 1 0.003411770      CPU         10      Addition     1
## 2 0.006412983      CPU         10      Addition     2
## 3 0.003450394      CPU         10      Addition     3
## 4 0.003098965      CPU         10      Addition     4
## 5 0.002490997      CPU         10      Addition     5
## 6 0.002594948      CPU         20      Addition     1
```

We tested three types of processors CPU, GPU, and TPU for three kinds of matrix operation, addition, multiplication, and inversion, with the matrix from size 10 to size 2160. We repeat each test for five times. **We measured $\log_{10}(\text{run-time})$ for each trial, and we use that as the evaluation of the performances.**

Simple Plots

Here is the general visualization for the performances of each processor under three matrix operations:

```
jpeg(filename = "../figs/overview.jpeg", width = 1000, height = 800, quality = 10000)
ggplot(data = data, aes(x = Processor, y = log10(Runtime))) +
  geom_boxplot(aes(fill = MatrixOperation))
while (!is.null(dev.list())) dev.off()
```



Zhanhao Zhang Part I

Pros & Cons of Each Processor

When matrix size is the same, is there any processor or operation effects? Or is there any interactive effect?

```
df_cpu <- data[data$Processor == "CPU",]
lm(Runtime ~ MatrixSize + as.factor(MatrixOperation), data = df_cpu) %>%
  summary()
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_cpu)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -234.810  -41.042   -6.483   48.307  248.752
##
## Coefficients:
##                                     Estimate Std. Error t value Pr(>|t|)
## (Intercept)                    -67.641331   13.699933  -4.937 2.37e-06
```

```
## MatrixSize                0.120877    0.009097  13.287 < 2e-16
## as.factor(MatrixOperation)Inversion    71.726373  17.944905   3.997 0.000107
## as.factor(MatrixOperation)Multiplication 55.821041  17.944905   3.111 0.002291
##
## (Intercept)                ***
## MatrixSize                  ***
## as.factor(MatrixOperation)Inversion    ***
## as.factor(MatrixOperation)Multiplication **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 85.12 on 131 degrees of freedom
## Multiple R-squared:  0.5971, Adjusted R-squared:  0.5879
## F-statistic: 64.73 on 3 and 131 DF,  p-value: < 2.2e-16
```

```
df_tpu <- data[data$Processor == "TPU",]
lm(Runtime ~ MatrixSize + as.factor(MatrixOperation), data = df_tpu) %>%
  summary()
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_tpu)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0008794 -0.0003316 -0.0002244 -0.0000993  0.0041166
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    5.460e-03  1.389e-04  39.296 <2e-16
## MatrixSize   -2.635e-08  9.226e-08  -0.286  0.776
## as.factor(MatrixOperation)Inversion -3.320e-03  1.820e-04 -18.243 <2e-16
## as.factor(MatrixOperation)Multiplication -2.238e-03  1.820e-04 -12.296 <2e-16
##
## (Intercept)                ***
## MatrixSize                  ***
## as.factor(MatrixOperation)Inversion    ***
## as.factor(MatrixOperation)Multiplication ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0008633 on 131 degrees of freedom
## Multiple R-squared:  0.7256, Adjusted R-squared:  0.7193
## F-statistic: 115.4 on 3 and 131 DF,  p-value: < 2.2e-16
```

```
df_gpu <- data[data$Processor == "GPU",]
lm(Runtime ~ MatrixSize + as.factor(MatrixOperation), data = df_gpu) %>%
  summary()
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixSize + as.factor(MatrixOperation),
```

```
##      data = df_gpu)
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -10.4178  -3.7532   0.9807    2.6766   24.7084
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   -2.9430365   1.0021808  -2.937  0.00392
## MatrixSize                     0.0051962   0.0006655    7.808 1.64e-12
## as.factor(MatrixOperation)Inversion    6.7906886   1.3127100    5.173 8.42e-07
## as.factor(MatrixOperation)Multiplication 0.0669200   1.3127100    0.051 0.95942
##
## (Intercept)                    **
## MatrixSize                     ***
## as.factor(MatrixOperation)Inversion    ***
## as.factor(MatrixOperation)Multiplication
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.227 on 131 degrees of freedom
## Multiple R-squared:  0.4237, Adjusted R-squared:  0.4105
## F-statistic: 32.1 on 3 and 131 DF, p-value: 1.273e-15
```

Lixian Chen Part I

Analysis

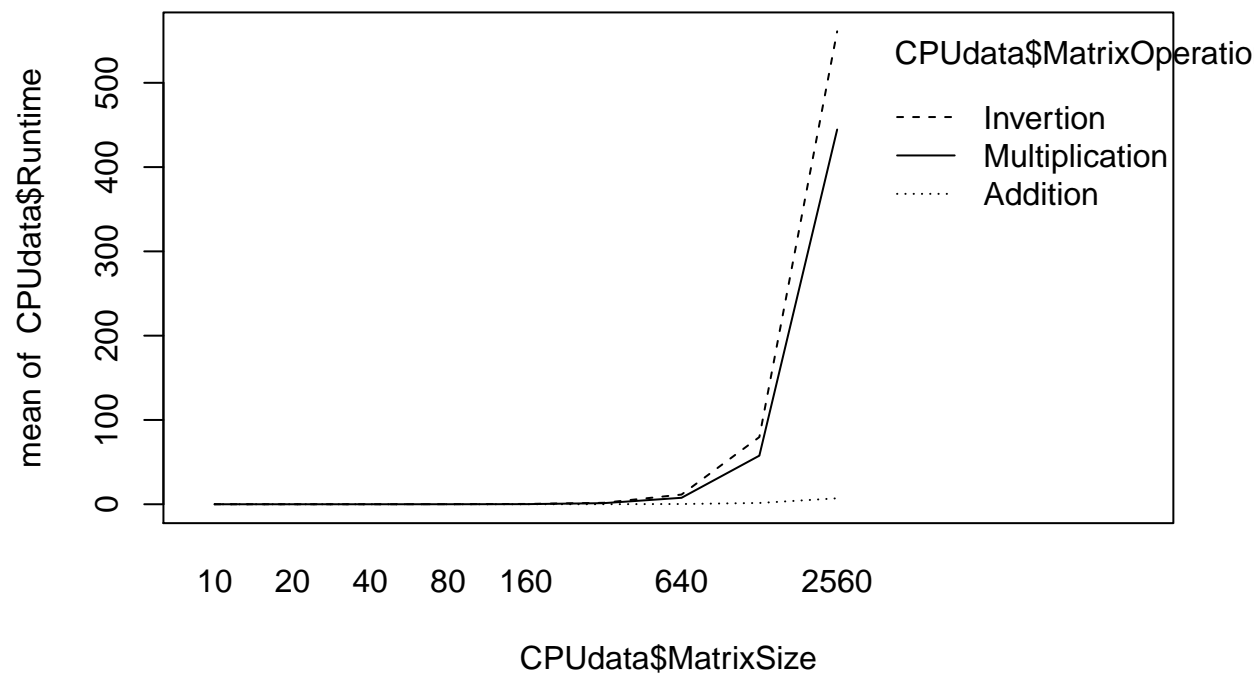
When the processor is the same, is there any operation and matrix size effect? We want to answer the following question: in each scenario, which processor should we use?

```
df <- fread("../data/Runtime.csv", header=TRUE)
attach(df)

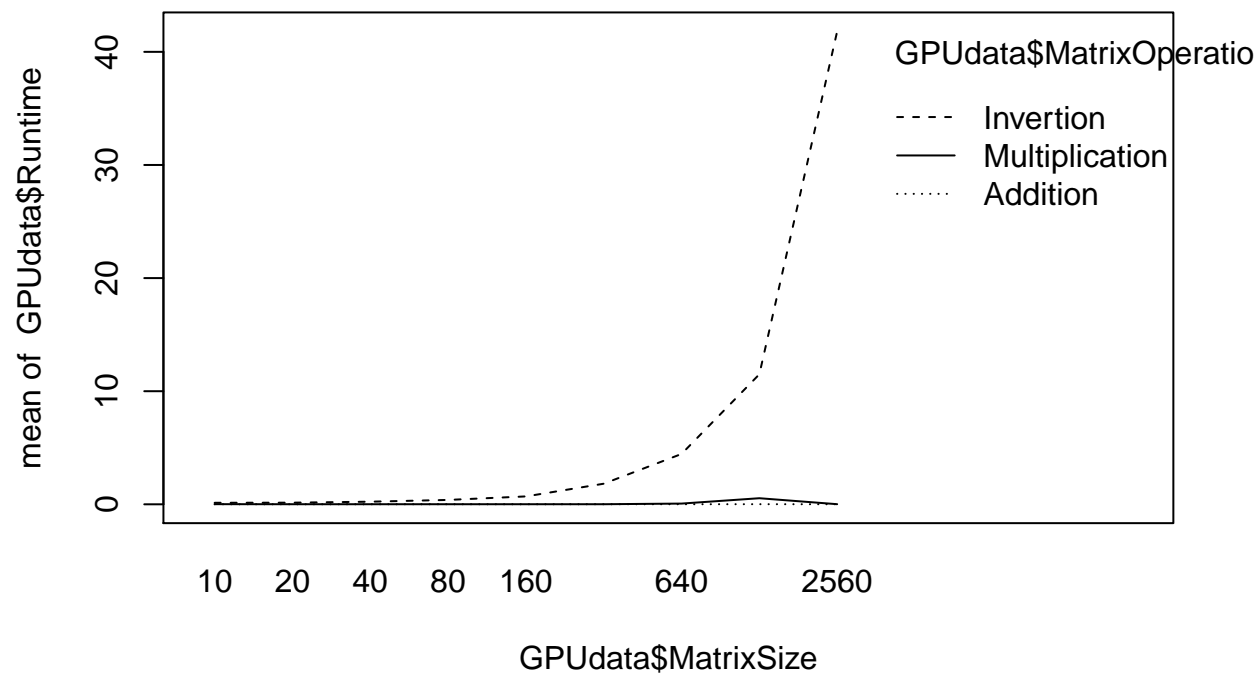
MatrixOperation<-factor(MatrixOperation)
Processor<-factor(Processor)

CPUdata <- df %>% filter(Processor=="CPU")
GPUdata <- df %>% filter(Processor=="GPU")
TPUdata <- df %>% filter(Processor=="TPU")

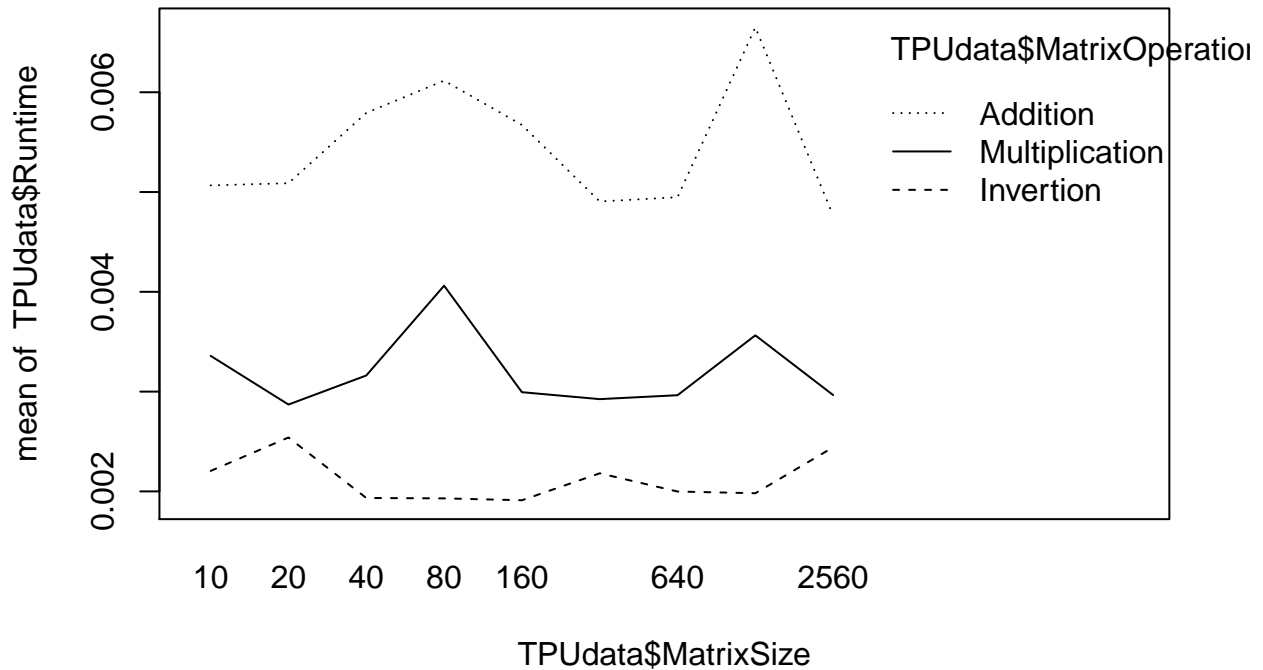
#op=par(mfrow=c(1,3))
interaction.plot(CPUdata$MatrixSize, CPUdata$MatrixOperation, CPUdata$Runtime)
```



```
interaction.plot(GPUdata$MatrixSize, GPUdata$MatrixOperation, GPUdata$Runtime)
```



```
interaction.plot(TPUdata$MatrixSize, TPUdata$MatrixOperation, TPUdata$Runtime)
```



```
#par(op)

library(psych)
describeBy(df$Runtime,group = df$MatrixSize, mat = TRUE) %>% #create dataframe
  select(MatrixSize=group1, N=n, Mean=mean, SD=sd, Median=median, Min=min, Max=max,
         Skew=skew, Kurtosis=kurtosis, SEM=se)
```

##	MatrixSize	N	Mean	SD	Median	Min	Max
##	X11	10 45	0.02117285	0.04190229	0.005617142	0.001893044	0.1475253
##	X12	20 45	0.02285517	0.04503445	0.004884481	0.001916409	0.1613362
##	X13	40 45	0.03386522	0.07211264	0.005309582	0.001902819	0.2506576
##	X14	80 45	0.06023479	0.12131721	0.006777525	0.001910210	0.4039948
##	X15	160 45	0.13455488	0.22988237	0.006753206	0.001870394	0.7055206
##	X16	320 45	0.54802946	0.77266887	0.008684158	0.001933098	1.8301501
##	X17	640 45	2.65963363	4.07807665	0.064900875	0.001965523	11.5108950
##	X18	1280 45	16.76243327	28.75264487	0.538181782	0.001962900	80.1931298
##	X19	2560 45	117.16431474	210.69591026	0.008469582	0.001975775	562.2822752
##	Skew	Kurtosis	SEM				
##	X11	2.2815817	3.52258988	0.006246425			
##	X12	2.2598210	3.47569251	0.006713340			
##	X13	2.3149204	3.63635579	0.010749918			
##	X14	2.1514343	3.03908564	0.018084903			
##	X15	1.6153802	1.18296414	0.034268841			
##	X16	0.7633051	-1.35170085	0.115182675			
##	X17	1.1600905	-0.21560402	0.607923773			

```
## X18 1.3579285 0.07252682 4.286191230
## X19 1.3393614 -0.10016629 31.408691862
```

```
#boxplot(Runtime~Processor*MatrixOperation)
#tapply(Runtime,list(Processor, MatrixOperation),mean)
#tapply(Runtime,MatrixOperation,mean)
```

```
group_by(df, Processor) %>%
  summarise(
    count = n(),
    mean = mean(Runtime, na.rm = TRUE),
    sd = sd(Runtime, na.rm = TRUE)
  )
```

```
## # A tibble: 3 x 4
##   Processor count      mean      sd
## * <chr>      <int>   <dbl>   <dbl>
## 1 CPU         135  43.5    133.
## 2 GPU         135   2.29     8.11
## 3 TPU         135   0.00359 0.00163
```

```
group_by(df, MatrixOperation) %>%
  summarise(
    count = n(),
    mean = mean(Runtime, na.rm = TRUE),
    sd = sd(Runtime, na.rm = TRUE)
  )
```

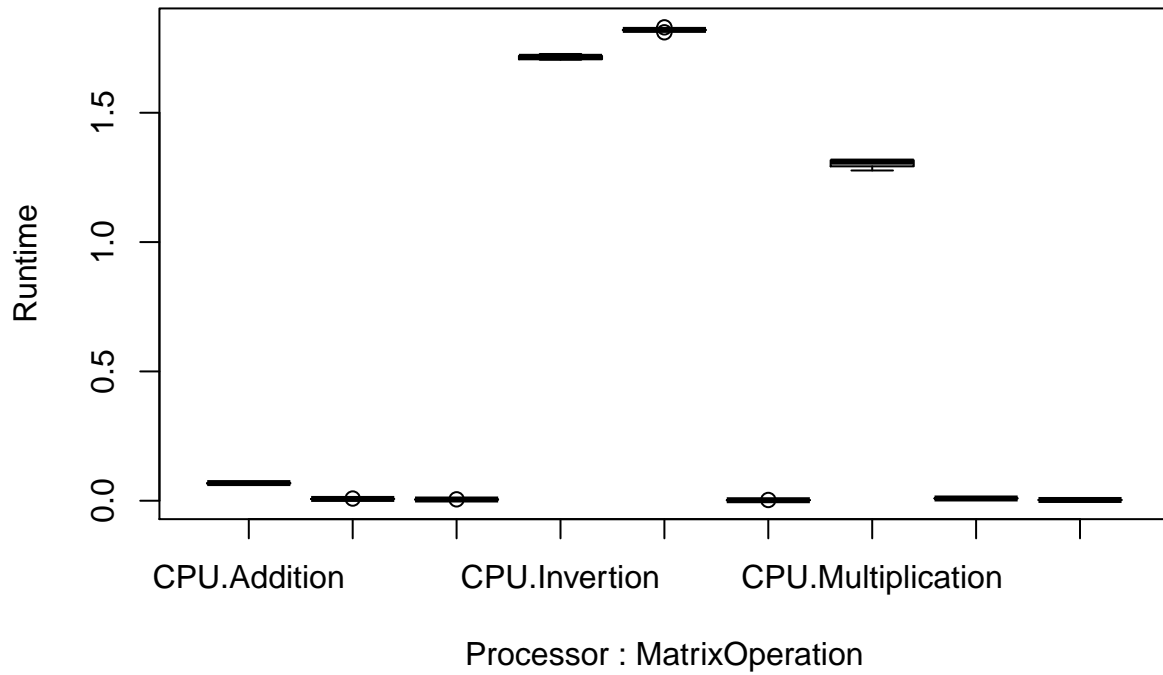
```
## # A tibble: 3 x 4
##   MatrixOperation count      mean      sd
## * <chr>          <int>   <dbl>   <dbl>
## 1 Addition        135   0.334    1.35
## 2 Inversion        135  26.5    107.
## 3 Multiplication   135  19.0    84.5
```

```
#detach(data)
```

```
size320data <- df %>% filter(MatrixSize==320)
size640data <- df %>% filter(MatrixSize==640)
size1280data <- df %>% filter(MatrixSize==1280)
size2560data <- df %>% filter(MatrixSize==2560)
```

```
boxplot(Runtime~Processor*MatrixOperation, data = size320data, main="At the level of matrix size=320")
```


At the level of matrix size=320

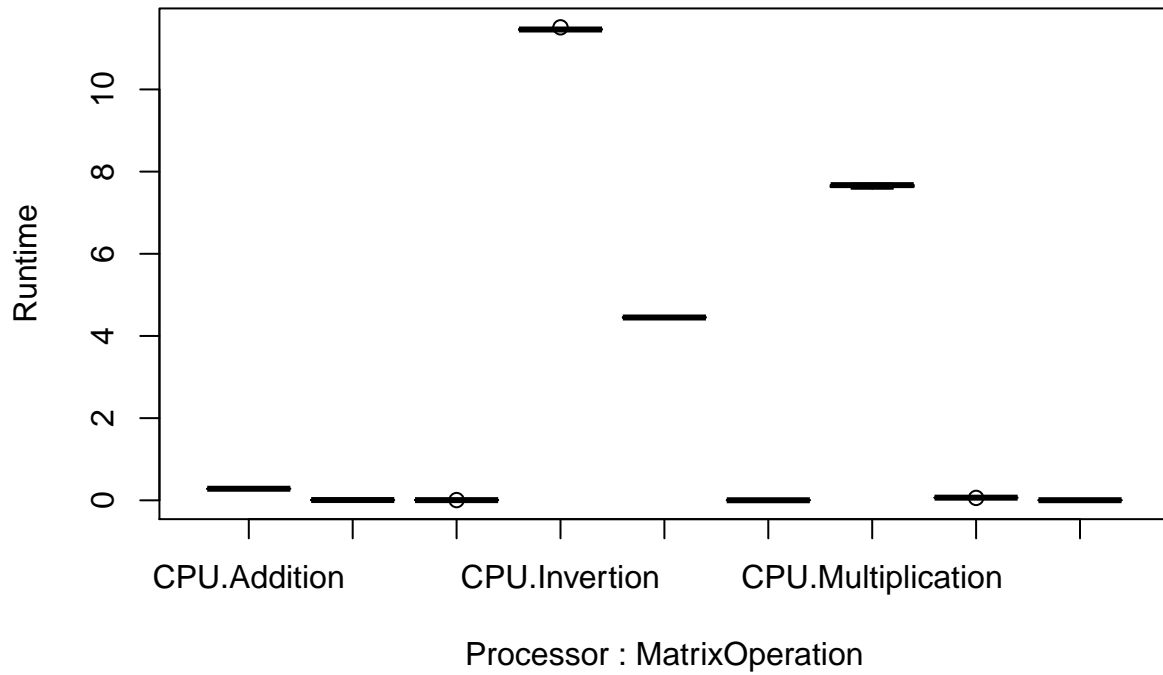


```
tapply(Runtime,list(Processor, MatrixOperation),mean, data = size320data)
```

```
##      Addition      Inversion Multiplication
## CPU 0.989863337 72.716236062 56.810904413
## GPU 0.007234912 6.797923491 0.074154954
## TPU 0.005444972 0.002124808 0.003207064
```

```
boxplot(Runtime~Processor*MatrixOperation, data = size640data, main="At the level of matrix size=640")
```

At the level of matrix size=640

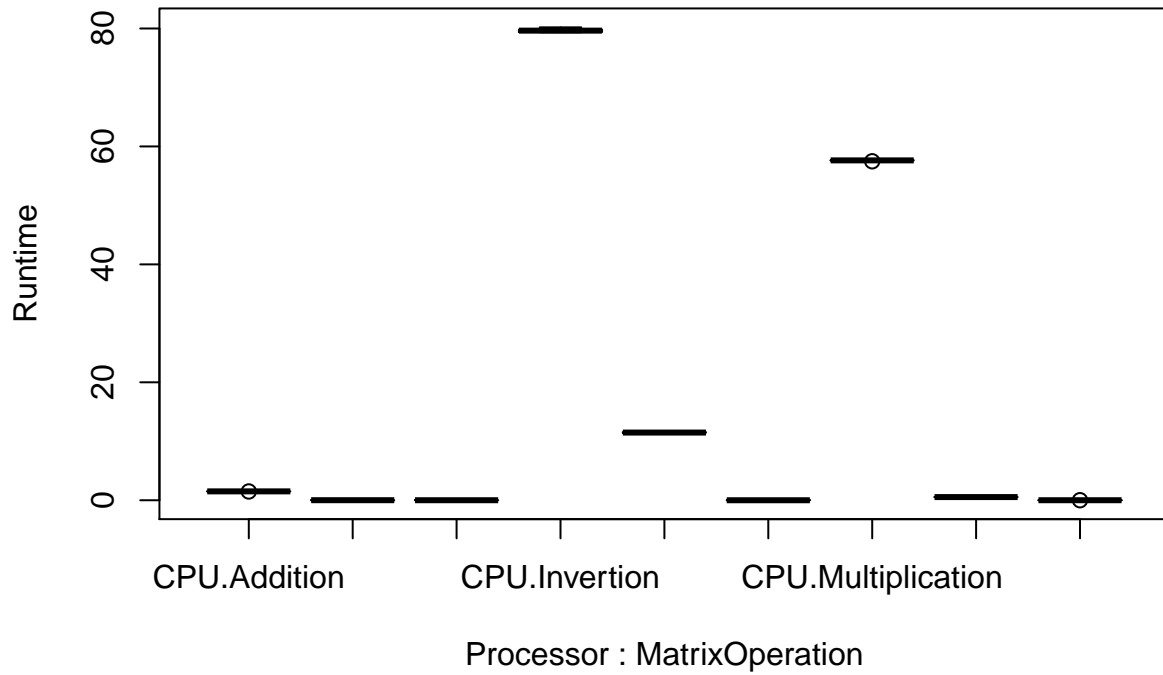


```
tapply(Runtime,list(Processor, MatrixOperation),mean, data = size640data)
```

```
##      Addition      Inversion Multiplication
## CPU 0.989863337 72.716236062 56.810904413
## GPU 0.007234912 6.797923491 0.074154954
## TPU 0.005444972 0.002124808 0.003207064
```

```
boxplot(Runtime~Processor*MatrixOperation, data = size1280data, main="At the level of matrix size=1280")
```

At the level of matrix size=1280

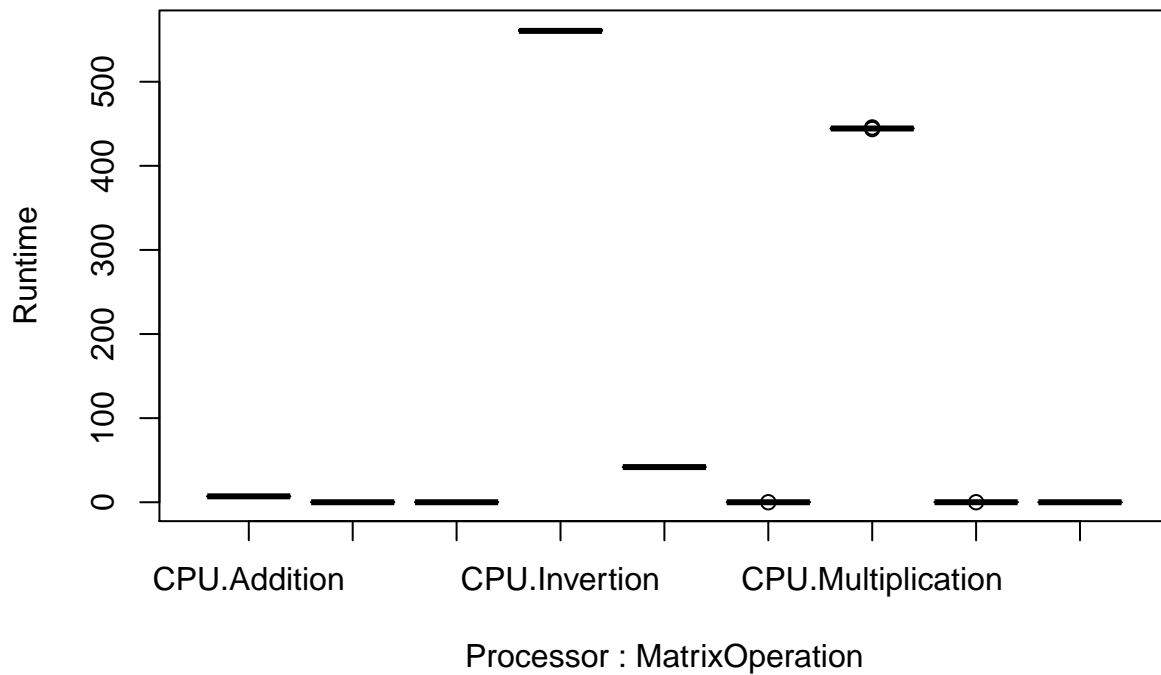


```
tapply(Runtime,list(Processor, MatrixOperation),mean, data = size1280data)
```

```
##      Addition      Inversion Multiplication
## CPU 0.989863337 72.716236062 56.810904413
## GPU 0.007234912 6.797923491 0.074154954
## TPU 0.005444972 0.002124808 0.003207064
```

```
boxplot(Runtime~Processor*MatrixOperation, data = size2560data, main="At the level of matrix size=2560")
```

At the level of matrix size=2560



```
tapply(Runtime,list(Processor, MatrixOperation),mean, data = size2560data)
```

```
##      Addition      Inversion Multiplication
## CPU 0.989863337 72.716236062 56.810904413
## GPU 0.007234912 6.797923491 0.074154954
## TPU 0.005444972 0.002124808 0.003207064
```

At the level of matrix size=2560, avoid using CPU for inversion and multiplication, because its runtime are much bigger

```
#320
```

```
fit2<-lm(Runtime~Processor+MatrixOperation, data = size320data)
summary(fit2)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.63252 -0.44075  0.05662  0.39766  0.58682
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.5072     0.1496   3.389  0.00159 **
```

```
## ProcessorGPU          -0.4163      0.1639  -2.540  0.01508 *
## ProcessorTPU          -1.0252      0.1639  -6.254  2.08e-07 ***
## MatrixOperationInversion  1.1525      0.1639   7.031  1.70e-08 ***
## MatrixOperationMultiplication 0.4116      0.1639   2.511  0.01619 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4489 on 40 degrees of freedom
## Multiple R-squared:  0.6931, Adjusted R-squared:  0.6624
## F-statistic: 22.59 on 4 and 40 DF,  p-value: 8.149e-10
```

```
fit1<-lm(Runtime~Processor*MatrixOperation, data = size320data)
summary(fit1)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0259142 -0.0004408 -0.0000341  0.0006995  0.0137136
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)      0.067873   0.003147   21.57
## ProcessorGPU     -0.060596   0.004450  -13.62
## ProcessorTPU     -0.062968   0.004450  -14.15
## MatrixOperationInversion  1.647123   0.004450  370.15
## MatrixOperationMultiplication 1.234854   0.004450  277.50
## ProcessorGPU:MatrixOperationInversion  0.165856   0.006293   26.36
## ProcessorTPU:MatrixOperationInversion -1.649847   0.006293 -262.17
## ProcessorGPU:MatrixOperationMultiplication -1.233006   0.006293 -195.93
## ProcessorTPU:MatrixOperationMultiplication -1.236834   0.006293 -196.54
##              Pr(>|t|)
## (Intercept)      < 2e-16 ***
## ProcessorGPU      9.04e-16 ***
## ProcessorTPU      2.80e-16 ***
## MatrixOperationInversion  < 2e-16 ***
## MatrixOperationMultiplication  < 2e-16 ***
## ProcessorGPU:MatrixOperationInversion  < 2e-16 ***
## ProcessorTPU:MatrixOperationInversion  < 2e-16 ***
## ProcessorGPU:MatrixOperationMultiplication  < 2e-16 ***
## ProcessorTPU:MatrixOperationMultiplication  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.007036 on 36 degrees of freedom
## Multiple R-squared:  0.9999, Adjusted R-squared:  0.9999
## F-statistic: 6.633e+04 on 8 and 36 DF,  p-value: < 2.2e-16
```

```
anova(fit2, fit1)
```

```
## Analysis of Variance Table
```

```
##
## Model 1: Runtime ~ Processor + MatrixOperation
## Model 2: Runtime ~ Processor * MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      40 8.0610
## 2      36 0.0018  4      8.0592 40700 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod320<-aov(Runtime~Processor*MatrixOperation, data = size320data)
#summary(mod320)
Anova(mod320,type="III")
```

```
## Anova Table (Type III tests)
##
## Response: Runtime
##
##               Sum Sq Df F value    Pr(>F)
## (Intercept)    0.0230  1   465.30 < 2.2e-16 ***
## Processor      0.0127  2    128.65 < 2.2e-16 ***
## MatrixOperation  7.3464  2 74200.69 < 2.2e-16 ***
## Processor:MatrixOperation 8.0592  4 40700.27 < 2.2e-16 ***
## Residuals      0.0018 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
modF<-lm(Runtime~Processor+MatrixOperation, data = size320data)
modA<-lm(Runtime~Processor, data = size320data)
modB<-lm(Runtime~MatrixOperation, data = size320data)

anova(modA, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 18.293
## 2      40  8.061  2     10.232 25.387 7.62e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(modB, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ MatrixOperation
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 16.036
## 2      40  8.061  2      7.9754 19.788 1.061e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

#640

```
fit2_640<-lm(Runtime~Processor+MatrixOperation, data = size640data)
summary(fit2)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.63252 -0.44075  0.05662  0.39766  0.58682
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.5072     0.1496   3.389  0.00159 **
## ProcessorGPU     -0.4163     0.1639  -2.540  0.01508 *
## ProcessorTPU     -1.0252     0.1639  -6.254 2.08e-07 ***
## MatrixOperationInversion  1.1525     0.1639   7.031 1.70e-08 ***
## MatrixOperationMultiplication  0.4116     0.1639   2.511  0.01619 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4489 on 40 degrees of freedom
## Multiple R-squared:  0.6931, Adjusted R-squared:  0.6624
## F-statistic: 22.59 on 4 and 40 DF,  p-value: 8.149e-10
```

```
fit1_640<-lm(Runtime~Processor*MatrixOperation, data = size640data)
summary(fit1_640)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation, data = size640data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.070744 -0.001994 -0.000012  0.000669  0.055012
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)      0.282343   0.008379   33.70
## ProcessorGPU     -0.275710   0.011849  -23.27
## ProcessorTPU     -0.277395   0.011849  -23.41
## MatrixOperationInversion  11.186470   0.011849  944.06
## MatrixOperationMultiplication  7.371087   0.011849  622.07
## ProcessorGPU:MatrixOperationInversion  -6.741354   0.016758 -402.29
## ProcessorTPU:MatrixOperationInversion -11.189419   0.016758 -667.73
## ProcessorGPU:MatrixOperationMultiplication -7.313893   0.016758 -436.45
## ProcessorTPU:MatrixOperationMultiplication -7.373071   0.016758 -439.99
##              Pr(>|t|)
## (Intercept)      <2e-16 ***
## ProcessorGPU     <2e-16 ***
## ProcessorTPU     <2e-16 ***
```

```
## MatrixOperationInversion          <2e-16 ***
## MatrixOperationMultiplication     <2e-16 ***
## ProcessorGPU:MatrixOperationInversion <2e-16 ***
## ProcessorTPU:MatrixOperationInversion <2e-16 ***
## ProcessorGPU:MatrixOperationMultiplication <2e-16 ***
## ProcessorTPU:MatrixOperationMultiplication <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01874 on 36 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 2.606e+05 on 8 and 36 DF, p-value: < 2.2e-16
```

```
anova(fit2_640, fit1_640)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor + MatrixOperation
## Model 2: Runtime ~ Processor * MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      40 184.706
## 2      36   0.013  4    184.69 131541 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod640<-aov(Runtime~Processor*MatrixOperation, data = size640data)
Anova(mod640,type="III")
```

```
## Anova Table (Type III tests)
##
## Response: Runtime
##
##           Sum Sq Df  F value    Pr(>F)
## (Intercept)    0.40  1  1135.52 < 2.2e-16 ***
## Processor      0.25  2   363.15 < 2.2e-16 ***
## MatrixOperation 323.38  2 460630.58 < 2.2e-16 ***
## Processor:MatrixOperation 184.69  4 131541.31 < 2.2e-16 ***
## Residuals      0.01 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
modF<-lm(Runtime~Processor+MatrixOperation, data = size640data)
modA<-lm(Runtime~Processor, data = size640data)
modB<-lm(Runtime~MatrixOperation, data = size640data)

anova(modA, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 388.42
```



```
## 2      40 184.71  2      203.71 22.058 3.496e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(modB, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ MatrixOperation
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 528.04
## 2      40 184.71  2      343.33 37.176 7.521e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#1280
```

```
fit2_1280<-lm(Runtime~Processor+MatrixOperation, data = size1280data)
summary(fit2)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.63252 -0.44075  0.05662  0.39766  0.58682
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.5072     0.1496   3.389  0.00159 **
## ProcessorGPU     -0.4163     0.1639  -2.540  0.01508 *
## ProcessorTPU     -1.0252     0.1639  -6.254 2.08e-07 ***
## MatrixOperationInversion  1.1525     0.1639   7.031 1.70e-08 ***
## MatrixOperationMultiplication  0.4116     0.1639   2.511  0.01619 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4489 on 40 degrees of freedom
## Multiple R-squared:  0.6931, Adjusted R-squared:  0.6624
## F-statistic: 22.59 on 4 and 40 DF,  p-value: 8.149e-10
```

```
fit1_1280<-lm(Runtime~Processor*MatrixOperation, data = size1280data)
summary(fit1_1280)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation, data = size1280data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.32048 -0.00256 -0.00001  0.00333  0.51330
```

```
##
## Coefficients:
##
##               Estimate Std. Error t value
## (Intercept)      1.50458    0.05293   28.43
## ProcessorGPU     -1.49482    0.07485  -19.97
## ProcessorTPU     -1.49793    0.07485  -20.01
## MatrixOperationInversion  78.17525    0.07485 1044.41
## MatrixOperationMultiplication  56.11665    0.07485   749.71
## ProcessorGPU:MatrixOperationInversion -66.68873    0.10586 -630.00
## ProcessorTPU:MatrixOperationInversion -78.17992    0.10586 -738.56
## ProcessorGPU:MatrixOperationMultiplication -55.58840    0.10586 -525.14
## ProcessorTPU:MatrixOperationMultiplication -56.11973    0.10586 -530.16
##
##               Pr(>|t|)
## (Intercept)      <2e-16 ***
## ProcessorGPU     <2e-16 ***
## ProcessorTPU     <2e-16 ***
## MatrixOperationInversion  <2e-16 ***
## MatrixOperationMultiplication  <2e-16 ***
## ProcessorGPU:MatrixOperationInversion  <2e-16 ***
## ProcessorTPU:MatrixOperationInversion  <2e-16 ***
## ProcessorGPU:MatrixOperationMultiplication  <2e-16 ***
## ProcessorTPU:MatrixOperationMultiplication  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1183 on 36 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 3.246e+05 on 8 and 36 DF, p-value: < 2.2e-16
```

```
anova(fit2_1280, fit1_1280)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor + MatrixOperation
## Model 2: Runtime ~ Processor * MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      40 9812.3
## 2      36   0.5  4    9811.8 175129 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
modF<-lm(Runtime~Processor+MatrixOperation, data = size1280data)
modA<-lm(Runtime~Processor, data = size1280data)
modB<-lm(Runtime~MatrixOperation, data = size1280data)

anova(modA, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 16666.1
```

```
## 2      40  9812.3  2      6853.7 13.97 2.505e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(modB, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ MatrixOperation
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 29521.7
## 2      40  9812.3  2      19709 40.173 2.708e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod1280<-aov(Runtime~Processor*MatrixOperation, data = size1280data)
Anova(mod1280,type="III")
```

```
## Anova Table (Type III tests)
##
## Response: Runtime
##
##              Sum Sq Df    F value    Pr(>F)
## (Intercept)      11.3  1    808.11 < 2.2e-16 ***
## Processor         7.5  2    266.44 < 2.2e-16 ***
## MatrixOperation 16245.1  2 579906.51 < 2.2e-16 ***
## Processor:MatrixOperation 9811.8  4 175128.60 < 2.2e-16 ***
## Residuals         0.5 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#2560
fit2_2560<-lm(Runtime~Processor+MatrixOperation, data = size2560data)
summary(fit2)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.63252 -0.44075  0.05662  0.39766  0.58682
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.5072     0.1496   3.389  0.00159 **
## ProcessorGPU     -0.4163     0.1639  -2.540  0.01508 *
## ProcessorTPU     -1.0252     0.1639  -6.254 2.08e-07 ***
## MatrixOperationInversion  1.1525     0.1639   7.031 1.70e-08 ***
## MatrixOperationMultiplication  0.4116     0.1639   2.511  0.01619 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.4489 on 40 degrees of freedom
## Multiple R-squared:  0.6931, Adjusted R-squared:  0.6624
## F-statistic: 22.59 on 4 and 40 DF,  p-value: 8.149e-10

fit1_2560<-lm(Runtime~Processor*MatrixOperation, data = size2560data)
summary(fit1_2560)

##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation, data = size2560data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.85301 -0.00118 -0.00007  0.00026  1.23586
##
## Coefficients:
##                                Estimate Std. Error t value
## (Intercept)                   7.0363    0.1765    39.87
## ProcessorGPU                  -7.0296    0.2496   -28.17
## ProcessorTPU                  -7.0315    0.2496   -28.18
## MatrixOperationInversion      554.0101    0.2496  2219.89
## MatrixOperationMultiplication  437.5216    0.2496  1753.13
## ProcessorGPU:MatrixOperationInversion -512.2048    0.3529 -1451.25
## ProcessorTPU:MatrixOperationInversion -554.0125    0.3529 -1569.71
## ProcessorGPU:MatrixOperationMultiplication -437.5188    0.3529 -1239.64
## ProcessorTPU:MatrixOperationMultiplication -437.5234    0.3529 -1239.65
##                                Pr(>|t|)
## (Intercept)                   <2e-16 ***
## ProcessorGPU                   <2e-16 ***
## ProcessorTPU                   <2e-16 ***
## MatrixOperationInversion       <2e-16 ***
## MatrixOperationMultiplication  <2e-16 ***
## ProcessorGPU:MatrixOperationInversion <2e-16 ***
## ProcessorTPU:MatrixOperationInversion <2e-16 ***
## ProcessorGPU:MatrixOperationMultiplication <2e-16 ***
## ProcessorTPU:MatrixOperationMultiplication <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3946 on 36 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 1.568e+06 on 8 and 36 DF,  p-value: < 2.2e-16
```

```
anova(fit2_2560, fit1_2560)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor + MatrixOperation
## Model 2: Runtime ~ Processor * MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      40 541548
## 2      36      6  4    541542 869482 < 2.2e-16 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

mod2560<-aov(Runtime~Processor*MatrixOperation, data = size2560data)
Anova(mod2560,type="III")
```

```
## Anova Table (Type III tests)
##
## Response: Runtime
##
##              Sum Sq Df    F value    Pr(>F)
## (Intercept)      248  1    1589.81 < 2.2e-16 ***
## Processor        165  2     529.08 < 2.2e-16 ***
## MatrixOperation 853203  2 2739748.40 < 2.2e-16 ***
## Processor:MatrixOperation 541542  4 869481.88 < 2.2e-16 ***
## Residuals          6 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
modF<-lm(Runtime~Processor+MatrixOperation, data = size2560data)
modA<-lm(Runtime~Processor, data = size2560data)
modB<-lm(Runtime~MatrixOperation, data = size2560data)

anova(modA, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 859034
## 2      40 541548  2   317486 11.725 9.829e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(modB, modF)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ MatrixOperation
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 1635796
## 2      40  541548  2   1094248 40.412 2.501e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Zhanhao Zhang Part II

General Visualization

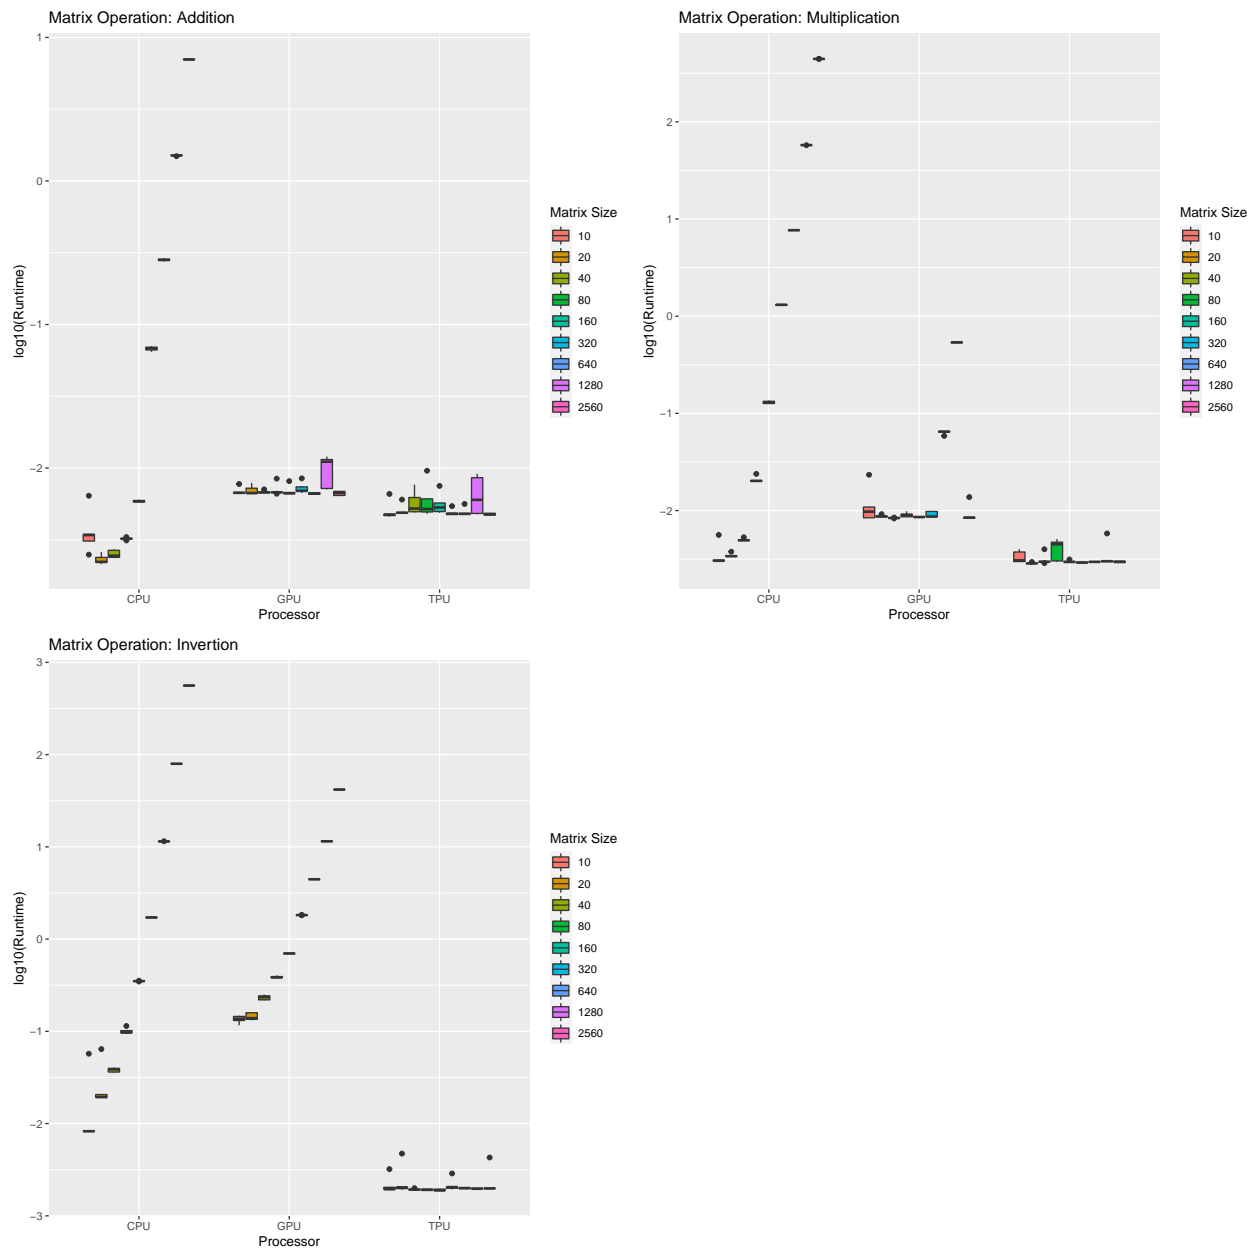


Figure 1: Operation v.s. Processors for Each Matrix Size

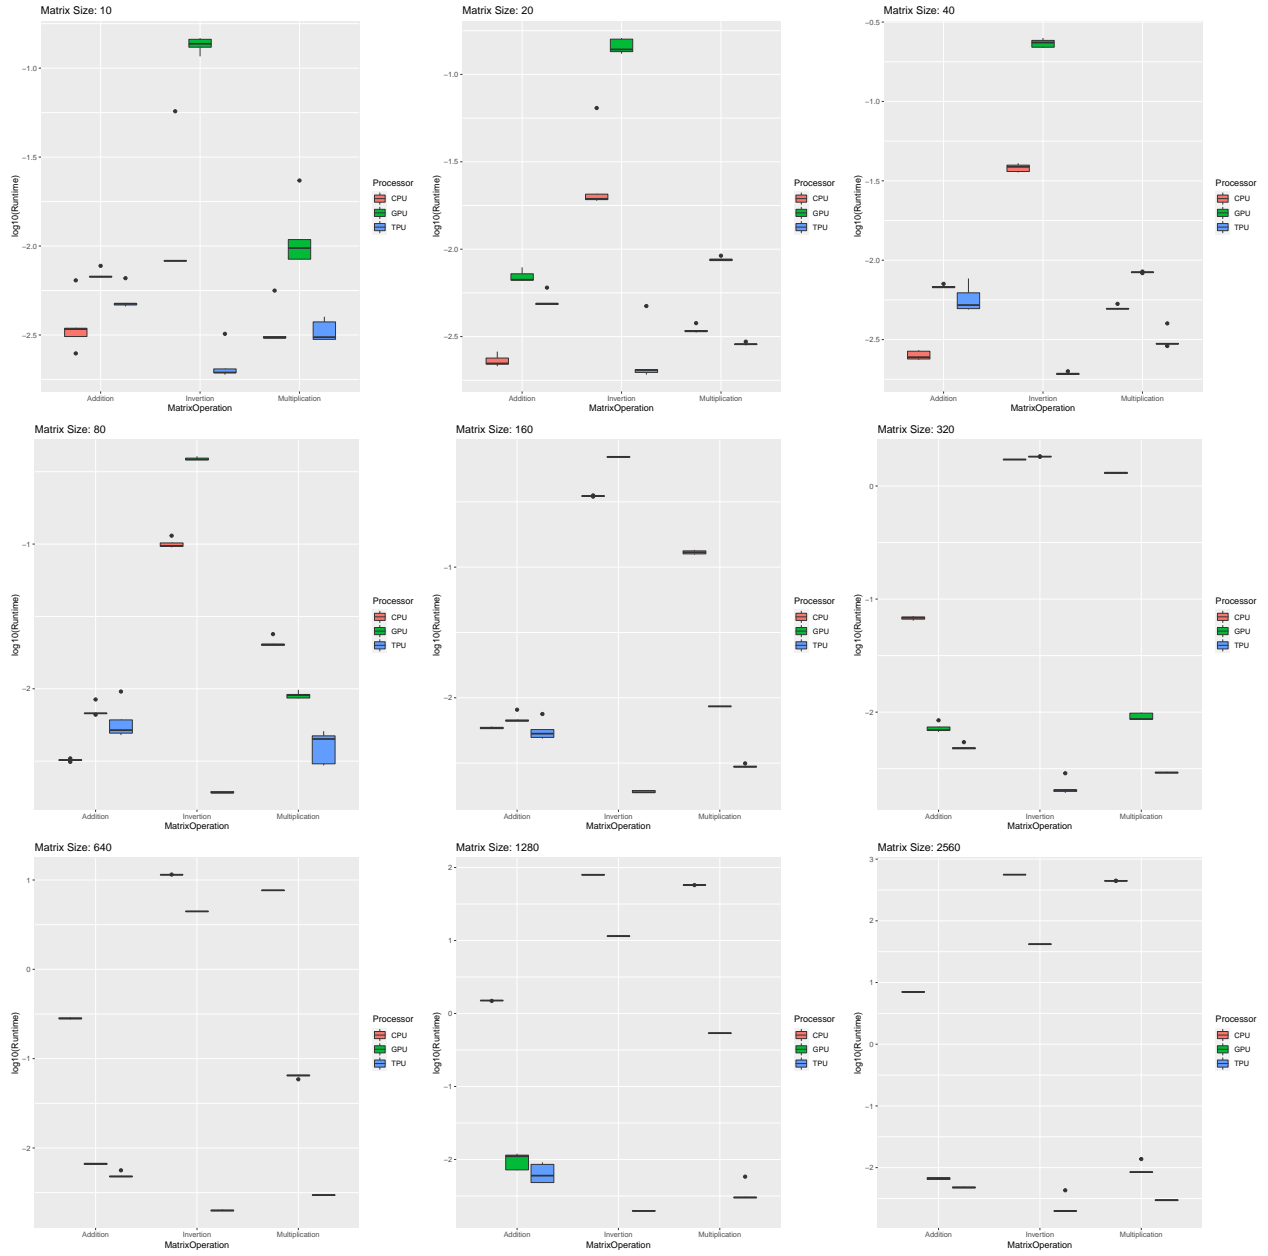


Figure 2: Matrix Size v.s. Operations for Each Processor.

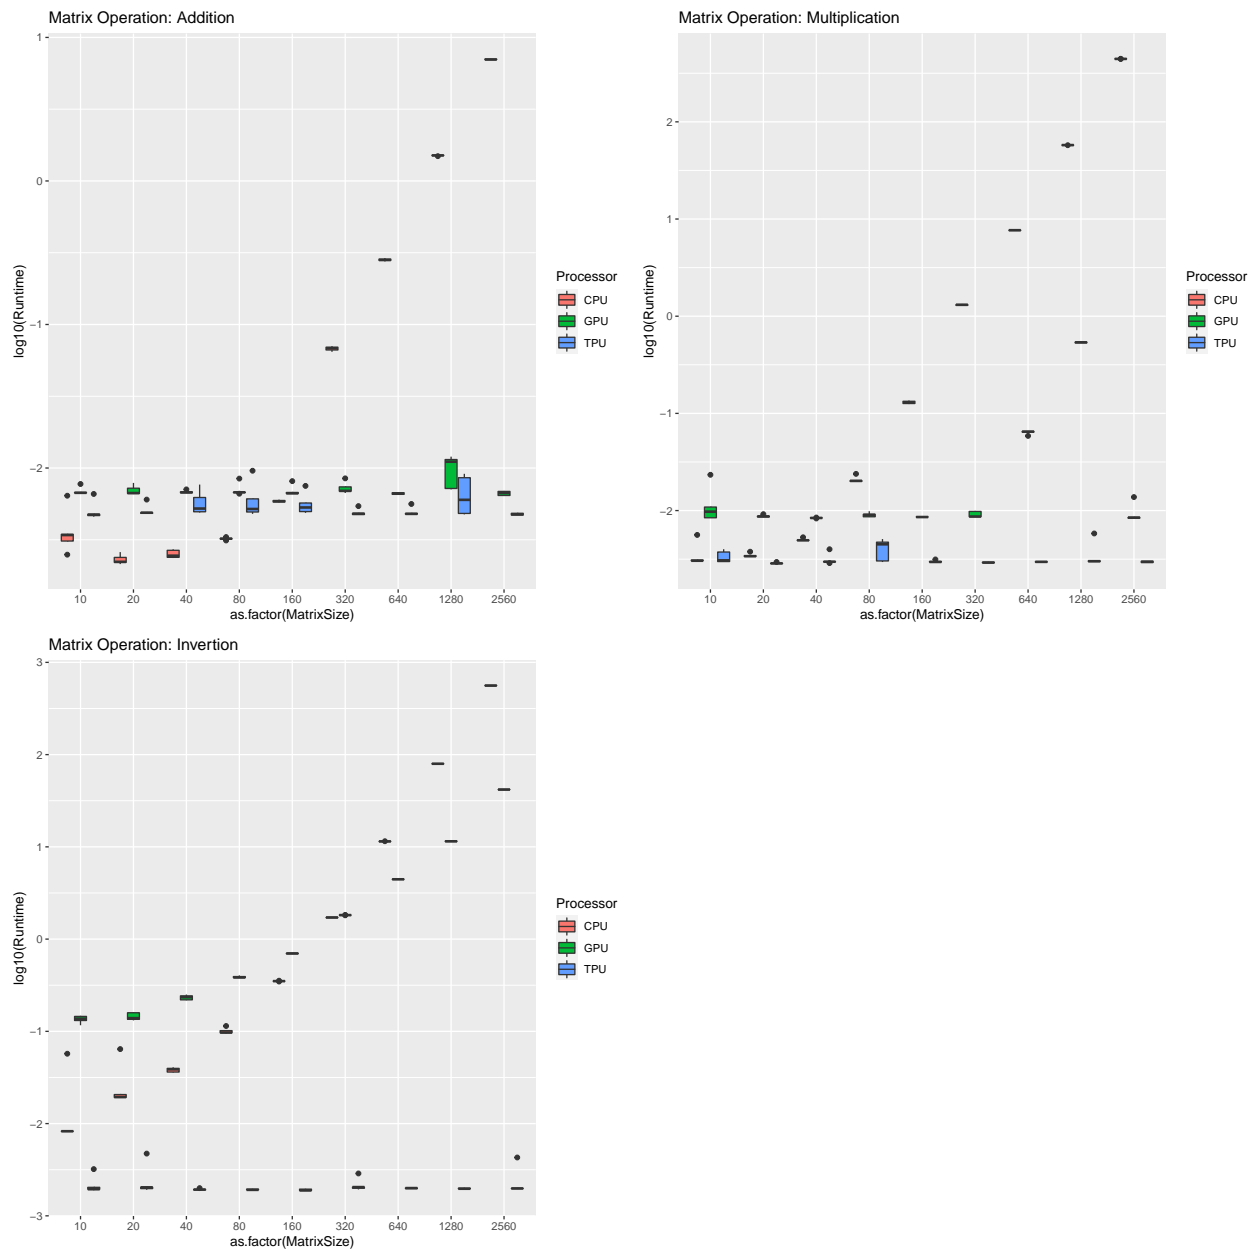



Figure 3: Operation v.s. Matrix Size for Each Processor

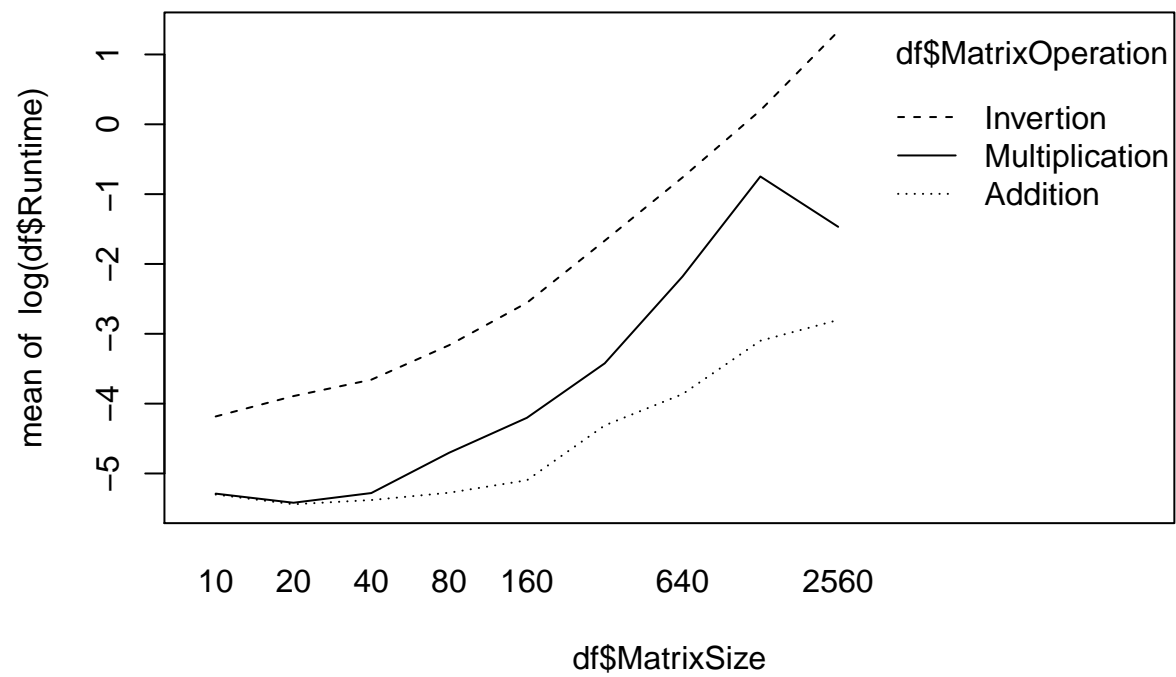
Lixian Chen Part II

Plots

```
#Interaction Plots  
op=par(mfrow=c(1,1))  
#op=par(mfrow=c(2,1))  
interaction.plot(df$MatrixSize, df$MatrixOperation, df$Runtime)
```



```
interaction.plot(df$MatrixSize, df$MatrixOperation, log(df$Runtime))
```



```
#par(op)
```