

Project on CPU, GPU, and TPU

Lixian Chen, Zhanhao Zhang, Jingbin Cao

4/22/2021

Preparing Required Packages

Read Data

```
data <- read.csv("../data/Runtime.csv")
head(data)
```


```
##      Runtime Processor MatrixSize MatrixOperation Trial
## 1 0.02059865      CPU          10      Addition      1
## 2 0.01871443      CPU          10      Addition      2
## 3 0.01863551      CPU          10      Addition      3
## 4 0.01847029      CPU          10      Addition      4
## 5 0.01866961      CPU          10      Addition      5
## 6 0.02328420      CPU          20      Addition      1
```

We tested three types of processors CPU, GPU, and TPU for three kinds of matrix operation, addition, multiplication, and inversion, with the matrix from size 10 to size 2160. We repeat each test for five times. **We measured $\log_{10}(\text{run-time})$ for each trial, and we use that as the evaluation of the performances.**

Simple Plots

Here is the general visualization for the performances of each processor under three matrix operations:

```
jpeg(filename = "../figs/overview.jpeg", width = 1000, height = 800, quality = 10000)
ggplot(data = data, aes(x = Processor, y = log10(Runtime))) +
  geom_boxplot(aes(fill = MatrixOperation))
while (!is.null(dev.list())) dev.off()
```



Jingbin Cao Part I

Research different Matrix Sizes One Way Anova for different matrix sized for each pair of processor and matrix operation: $\mu_1 = MatrixSize = 160$ $\mu_2 = MatrixSize = 320$ $\mu_3 = MatrixSize = 640$ $\mu_4 = MatrixSize = 1280$

Getting Data

```
cpu_add <- data[data$Processor == "CPU" & data$MatrixOperation=="Addition" & data$MatrixSize >= 160,]
cpu_mult <- data[data$Processor == "CPU" & data$MatrixOperation=="Multiplication" & data$MatrixSize >= 160,]
cpu_inv <- data[data$Processor == "CPU" & data$MatrixOperation=="Inversion" & data$MatrixSize >= 160,]
gpu_add <- data[data$Processor == "GPU" & data$MatrixOperation=="Addition" & data$MatrixSize >= 160,]
gpu_mult <- data[data$Processor == "GPU" & data$MatrixOperation=="Multiplication" & data$MatrixSize >= 160,]
gpu_inv <- data[data$Processor == "GPU" & data$MatrixOperation=="Inversion" & data$MatrixSize >= 160,]
tpu_add <- data[data$Processor == "TPU" & data$MatrixOperation=="Addition" & data$MatrixSize >= 160,]
tpu_mult <- data[data$Processor == "TPU" & data$MatrixOperation=="Multiplication" & data$MatrixSize >= 160,]
tpu_inv <- data[data$Processor == "TPU" & data$MatrixOperation=="Inversion" & data$MatrixSize >= 160,]
```

Anovas

```
summary(mod_cpu_add <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=cpu_add))
```

```
##               Df Sum Sq Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3 15.478    5.159   76040 <2e-16 ***
## Residuals              16  0.001    0.000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(mod_cpu_mult <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=cpu_mult))
```

```
##               Df Sum Sq Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3  19.53    6.511  393855 <2e-16 ***
## Residuals              16  0.00    0.000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(mod_cpu_inv <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=cpu_inv))
```

```
##               Df Sum Sq Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3  15.67    5.223  653718 <2e-16 ***
## Residuals              16  0.00    0.000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(mod_gpu_add <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=gpu_add))
```

```
##               Df Sum Sq Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3  2.6297  0.8766   10205 <2e-16 ***
## Residuals              16 0.0014  0.0001
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(mod_gpu_mult <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=gpu_mult))
```

```
##               Df Sum Sq Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3 12.328    4.109   34853 <2e-16 ***
## Residuals              16  0.002    0.000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(mod_gpu_inv <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=gpu_inv))
```

```
##               Df Sum Sq Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3  4.083    1.361  4518902 <2e-16 ***
## Residuals              16  0.000    0.000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(mod_tpu_add <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=tpu_add))
```

```
##                Df    Sum Sq  Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3 0.003245 0.001081   0.849  0.487
## Residuals              16 0.020376 0.001273
```

```
summary(mod_tpu_mult <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=tpu_mult))
```

```
##                Df    Sum Sq  Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3 0.00678 0.002259   0.707  0.562
## Residuals              16 0.05116 0.003197
```

```
summary(mod_tpu_inv <- aov(log10(Runtime) ~ as.factor(MatrixSize), data=tpu_inv))
```

```
##                Df    Sum Sq  Mean Sq F value Pr(>F)
## as.factor(MatrixSize)  3 0.00667 0.002224   0.936  0.446
## Residuals              16 0.03801 0.002376
```

```
# From the tables, we can see that only TPU & Multiplication and TPU & Inversion do not have Matrix Size
# plot(aov(Runtime ~ as.factor(MatrixSize), data=tpu_mult))
# plot(aov(Runtime ~ as.factor(MatrixSize), data=tpu_inv))
```

Confidence Interval for LMs

```
round(digits=4, confint(mod_cpu_add))
```

```
##                2.5 % 97.5 %
## (Intercept)      -1.1994 -1.1838
## as.factor(MatrixSize)320    1.0339 1.0560
## as.factor(MatrixSize)640    1.6346 1.6567
## as.factor(MatrixSize)1280   2.3952 2.4173
```

```
round(digits=4, confint(mod_cpu_mult))
```

```
##                2.5 % 97.5 %
## (Intercept)      0.1445 0.1522
## as.factor(MatrixSize)320  0.9245 0.9354
## as.factor(MatrixSize)640  1.7674 1.7783
## as.factor(MatrixSize)1280 2.6596 2.6705
```

```
round(digits=4, confint(mod_cpu_inv))
```

```
##                2.5 % 97.5 %
## (Intercept)      0.5928 0.5982
## as.factor(MatrixSize)320  0.6918 0.6994
## as.factor(MatrixSize)640  1.5130 1.5206
## as.factor(MatrixSize)1280 2.3589 2.3664
```

```
round(confint(mod_gpu_add),4)
```

```
##                2.5 % 97.5 %
## (Intercept)    -1.1665 -1.1489
## as.factor(MatrixSize)320 -0.0066 0.0183
## as.factor(MatrixSize)640 0.0865 0.1114
## as.factor(MatrixSize)1280 0.8550 0.8798
```

```
round(digits=4,confint(mod_gpu_mult))
```

```
##                2.5 % 97.5 %
## (Intercept)    -1.0610 -1.0404
## as.factor(MatrixSize)320 0.3907 0.4198
## as.factor(MatrixSize)640 1.1183 1.1474
## as.factor(MatrixSize)1280 2.0502 2.0793
```


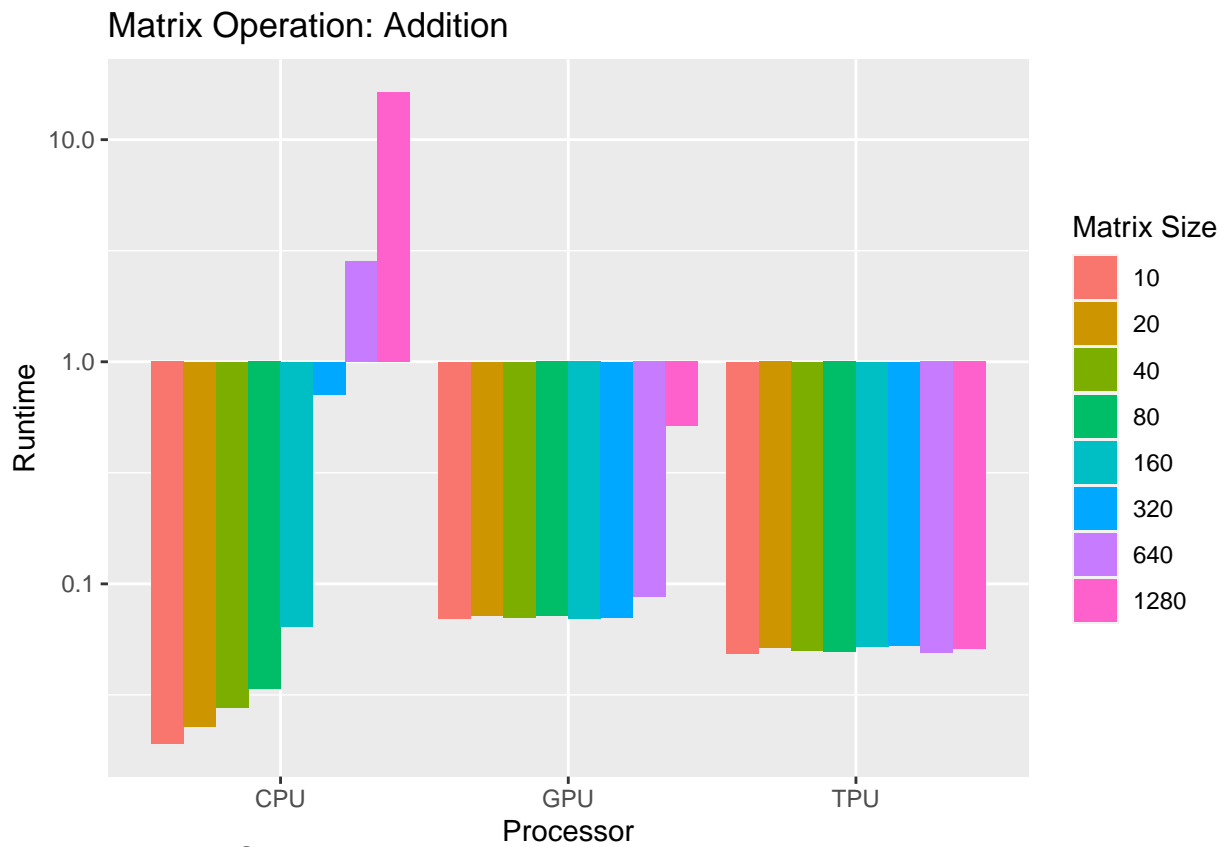
```
round(digits=4,confint(mod_gpu_inv))
```


```
##                2.5 % 97.5 %
## (Intercept)    0.8389 0.8400
## as.factor(MatrixSize)320 0.4203 0.4218
## as.factor(MatrixSize)640 0.8101 0.8116
## as.factor(MatrixSize)1280 1.2162 1.2177
```

Zhanhao Zhang Part I



Interaction Plots



```
#ggplot(data = data, aes(x = Processor, y = log10(Runtime))) +
# geom_boxplot(aes(fill = MatrixOperation))
for(operation in unique(data$MatrixOperation)){
  #png(paste0("../figs/", operation, ".png"), width = 500, height = 500)
  p <- data %>%
    dplyr::filter(MatrixOperation == operation) %>%
    group_by(MatrixSize, MatrixOperation, Processor) %>%
    summarize(Runtime = mean(Runtime)) %>%
    ggplot(aes(x = Processor, y = Runtime)) +
    #geom_boxplot(aes(fill = as.factor(MatrixSize))) +
    geom_bar(aes(fill = as.factor(MatrixSize)), stat = "identity",
              position = "dodge") +
    scale_y_log10() +
    ggtitle(paste0("Matrix Operation: ", operation)) +
    #facet_wrap(~ MatrixOperation, scales = "free", nrow = 1) +
    guides(fill=guide_legend(title = "Matrix Size"))
  print(p)
  #dev.off()
}
```






```
for(size in unique(data$MatrixSize)){
  p <- data %>%
    dplyr::filter(MatrixSize == size) %>%
    group_by(MatrixSize, MatrixOperation, Processor) %>%
    summarize(Runtime = mean(Runtime)) %>%
    ggplot(aes(x = Processor, y = Runtime)) +
    #geom_boxplot(aes(fill = as.factor(MatrixOperation))) +
    geom_bar(aes(fill = as.factor(MatrixOperation)), position = "dodge",
              stat = "identity") +
    scale_y_log10() +
    ggtitle(paste0("Matrix Size: ", size)) +
    #facet_wrap(~ MatrixOperation, scales = "free", nrow = 1) +
    guides(fill=guide_legend(title = "Matrix Operation"))
  print(p)
}
```








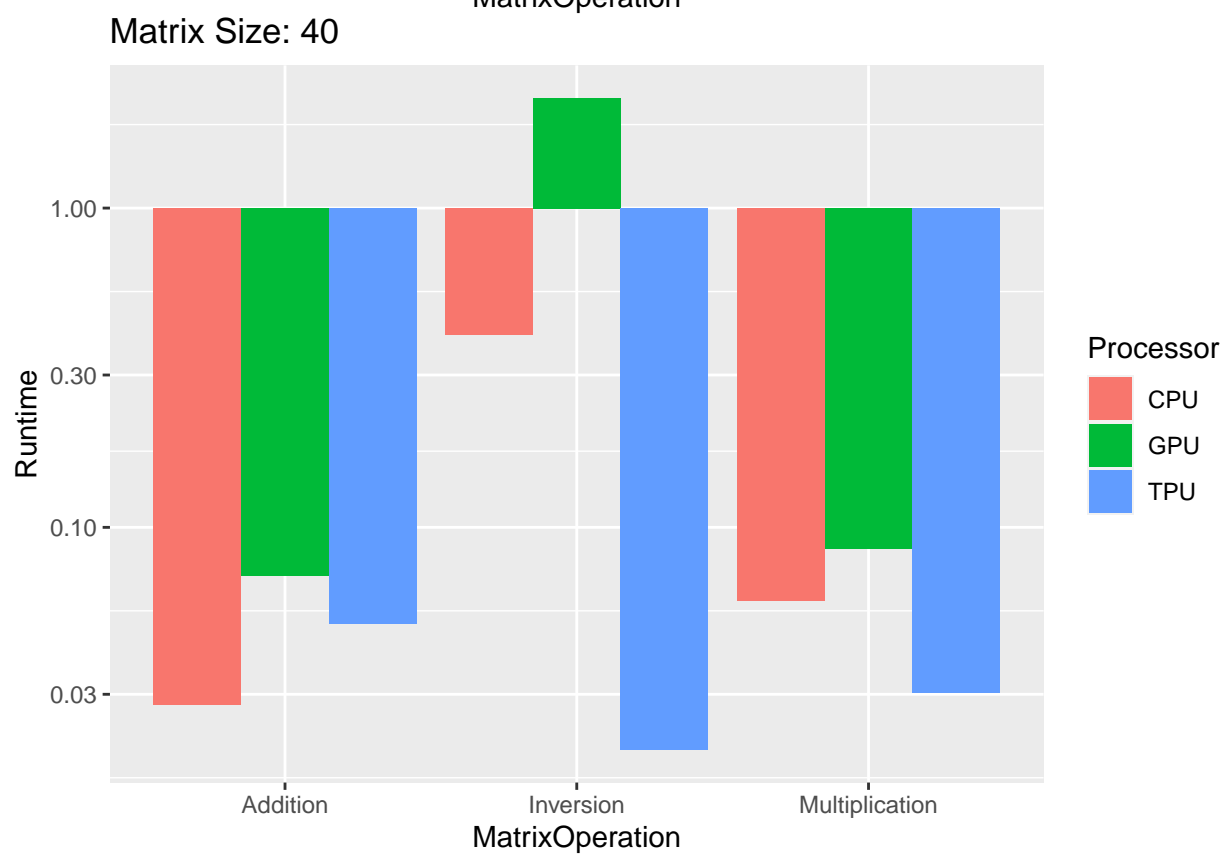





```



for(size in unique(data$MatrixSize)){
  p <- data %>%
    dplyr::filter(MatrixSize == size) %>%
    group_by(MatrixSize, MatrixOperation, Processor) %>%
    summarize(Runtime = mean(Runtime)) %>%
    ggplot(aes(x = MatrixOperation, y = Runtime)) +
    #geom_boxplot(aes(fill = as.factor(Processor))) +
    geom_bar(aes(fill = as.factor(Processor)), position = "dodge",
             stat = "identity") +
    scale_y_log10() +
    ggtitle(paste0("Matrix Size: ", size)) +
    #facet_wrap(~ MatrixOperation, scales = "free", nrow = 1) +
    guides(fill=guide_legend(title = "Processor"))
  print(p)
}


```














```
for(operation in unique(data$MatrixOperation)){
  p <- data %>%
    dplyr::filter(MatrixOperation == operation) %>%
    group_by(MatrixSize, MatrixOperation, Processor) %>%
    summarize(Runtime = mean(Runtime)) %>%
    ggplot(aes(x = as.factor(MatrixSize), y = Runtime)) +
    #geom_boxplot(aes(fill = as.factor(Processor))) +
    geom_bar(aes(fill = as.factor(Processor)), stat = "identity",
             alpha = 1, position = "dodge") +
    scale_y_log10() +
    ggtitle(paste0("Matrix Operation: ", operation)) +
    #facet_wrap(~ MatrixOperation, scales = "free", nrow = 1) +
    guides(fill=guide_legend(title = "Processor"))
  print(p)
}
```



Pros & Cons of Each Processor


CPU:

```
df_cpu <- data[data$Processor == "CPU",]
linreg_cpu <- lm(Runtime ~ MatrixSize + as.factor(MatrixOperation),
                 data = df_cpu)
summary(linreg_cpu) %>% print()
```


```
##
## Call:
## lm(formula = Runtime ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_cpu)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -355.80  -71.41   -8.16   66.26  411.40
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -120.02352    24.18254  -4.963 2.40e-06
## MatrixSize      0.38443     0.03081  12.477 < 2e-16
## as.factor(MatrixOperation)Inversion    130.34056    31.25182   4.171 5.89e-05
## as.factor(MatrixOperation)Multiplication  90.96003    31.25182   2.911 0.00433
##
```

```
## (Intercept) ***
## MatrixSize ***
## as.factor(MatrixOperation)Inversion ***
## as.factor(MatrixOperation)Multiplication **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 139.8 on 116 degrees of freedom
## Multiple R-squared:  0.6, Adjusted R-squared:  0.5896
## F-statistic: 57.99 on 3 and 116 DF, p-value: < 2.2e-16
```

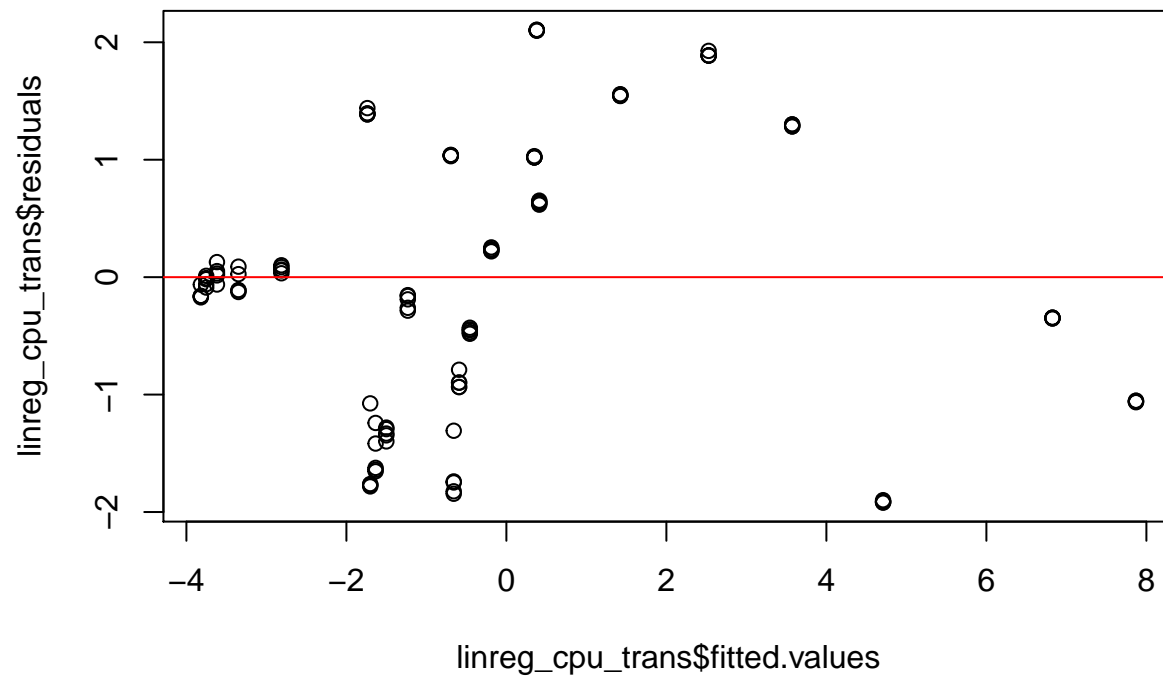
```
plot(linreg_cpu$residuals)
abline(h=0, col="red")
```



```
boxCox(linreg_cpu)
```



```
linreg_cpu_trans <- lm(log(Runtime) ~ MatrixSize + as.factor(MatrixOperation),
  data = df_cpu)
plot(linreg_cpu_trans$fitted.values, linreg_cpu_trans$residuals)
abline(h=0, col="red")
```




GPU:


```
df_gpu <- data[(data$Processor == "GPU") & (data$MatrixOperation != "Multiplication"),]
linreg_gpu <- lm(Runtime ~ MatrixSize + as.factor(MatrixOperation),
  data = df_gpu)
summary(linreg_gpu) %>% print()
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_gpu)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -41.981  -9.566  -2.974   10.907   47.604
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -13.920265    3.377960  -4.121 9.43e-05 ***
## MatrixSize      0.044073    0.005066   8.699 4.57e-13 ***
## as.factor(MatrixOperation)Inversion  23.893954    4.195848   5.695 2.15e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 18.76 on 77 degrees of freedom
## Multiple R-squared:  0.584, Adjusted R-squared:  0.5732
## F-statistic: 54.05 on 2 and 77 DF, p-value: 2.161e-15
```

```
plot(linreg_gpu$residuals)
abline(h=0, col="red")
```



```
boxCox(linreg_gpu)
```



```
linreg_gpu_trans <- lm(log10(Runtime) ~ MatrixSize +
                        as.factor(MatrixOperation), data = df_gpu)
summary(linreg_gpu_trans)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_gpu)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-0.50239	-0.23761	0.09982	0.19255	0.44349

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.374e+00	5.133e-02	-26.77	<2e-16 ***
MatrixSize	1.075e-03	7.699e-05	13.96	<2e-16 ***
as.factor(MatrixOperation)Inversion	1.894e+00	6.376e-02	29.70	<2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2851 on 77 degrees of freedom
## Multiple R-squared:  0.9333, Adjusted R-squared:  0.9315
## F-statistic: 538.4 on 2 and 77 DF,  p-value: < 2.2e-16
```

```
plot(linreg_gpu_trans$residuals)
abline(h=0, col="red")
```




TPU:

```
df_tpu <- data[data$Processor == "TPU",]
linreg_tpu <- lm(Runtime ~ MatrixSize + as.factor(MatrixOperation),
  data = df_tpu)
summary(linreg_tpu) %>% print()
```


```
##
## Call:
## lm(formula = Runtime ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_tpu)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0034224 -0.0020056 -0.0010012 -0.0000435  0.0270875
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    5.039e-02  7.237e-04   69.62  <2e-16
## MatrixSize     3.595e-07  9.221e-07    0.39   0.697
## as.factor(MatrixOperation)Inversion  -3.058e-02  9.353e-04  -32.70  <2e-16
## as.factor(MatrixOperation)Multiplication -1.877e-02  9.353e-04  -20.07  <2e-16
##
## (Intercept)          ***
## MatrixSize
## as.factor(MatrixOperation)Inversion      ***
## as.factor(MatrixOperation)Multiplication ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.004183 on 116 degrees of freedom
## Multiple R-squared:  0.9036, Adjusted R-squared:  0.9012
```

```
## F-statistic: 362.6 on 3 and 116 DF,  p-value: < 2.2e-16
```

```
plot(linreg_tpu$residuals)
abline(h=0, col="red")
```



```
boxCox(linreg_tpu)
```




```
linreg_tpu_trans <- lm(log10(Runtime) ~ MatrixSize +
                        as.factor(MatrixOperation), data = df_tpu)
summary(linreg_tpu_trans) %>% print()
```




```
##
## Call:
## lm(formula = log10(Runtime) ~ MatrixSize + as.factor(MatrixOperation),
##     data = df_tpu)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.043314 -0.023467 -0.014163  0.003178  0.273842
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)    -1.299e+00  8.067e-03 -161.041
## MatrixSize      4.439e-06  1.028e-05   0.432
## as.factor(MatrixOperation)Inversion  -4.050e-01  1.043e-02  -38.852
## as.factor(MatrixOperation)Multiplication -2.061e-01  1.043e-02  -19.770
##              Pr(>|t|)
## (Intercept)    <2e-16 ***
## MatrixSize      0.667
## as.factor(MatrixOperation)Inversion    <2e-16 ***
## as.factor(MatrixOperation)Multiplication <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04662 on 116 degrees of freedom
## Multiple R-squared:  0.9287, Adjusted R-squared:  0.9268
## F-statistic: 503.3 on 3 and 116 DF, p-value: < 2.2e-16
```

```
m <- max(abs(linreg_tpu_trans$residuals))
plot(linreg_tpu_trans$residuals, ylim=c(-m, m))
abline(h=0, col="red")
```




GPU Addition


```
df_small <- data[(data$Processor == "GPU") & (data$MatrixOperation == "Addition"),]
linreg_small <- lm(Runtime ~ MatrixSize, data = df_small)
plot(linreg_small$residuals)
abline(h=0, col="red")
```




```
boxCox(linreg_small, lambda = seq(-20, 2, 0.1))
```



```
linreg_small_transform <- lm(log10(Runtime) ~ MatrixSize, data = df_small)
plot(linreg_small_transform$residuals)
abline(h=0, col="red")
```




```
plot(df_small$MatrixSize, log10(df_small$Runtime))
abline(a=linreg_small_transform$coefficients[1],
       b=linreg_small_transform$coefficients[2])
```




GPU Multiplication

```
df_small <- data[(data$Processor == "GPU") & (data$MatrixOperation == "Multiplication"),]
linreg_small <- lm(Runtime ~ MatrixSize, data = df_small)
plot(linreg_small$residuals)
abline(h=0, col="red")
```



```
boxCox(linreg_small, lambda = seq(-10, 2, 0.1))
```




```
linreg_small_transform <- lm(log10(Runtime) ~ MatrixSize, data = df_small)
summary(linreg_small_transform)
```


```
##
## Call:
## lm(formula = log10(Runtime) ~ MatrixSize, data = df_small)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.189553 -0.042371 -0.006256  0.062402  0.136090
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.148e+00  1.822e-02  -63.01  <2e-16 ***
## MatrixSize   1.715e-03  3.487e-05   49.19  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09131 on 38 degrees of freedom
## Multiple R-squared:  0.9845, Adjusted R-squared:  0.9841
## F-statistic: 2419 on 1 and 38 DF,  p-value: < 2.2e-16
```

```
plot(linreg_small_transform$residuals)
abline(h=0, col="red")
```




```
plot(df_small$MatrixSize, log10(df_small$Runtime))
abline(a=linreg_small_transform$coefficients[1],
       b=linreg_small_transform$coefficients[2])
```




GPU Inversion

```
df_small <- data[(data$Processor == "GPU") & (data$MatrixOperation == "Inversion"),]
linreg_small <- lm(Runtime ~ MatrixSize, data = df_small)
plot(linreg_small$fitted.values, linreg_small$residuals)
abline(h=0, col="red")
```




```
boxCox(linreg_small, lambda = seq(-2, 2, 0.1))
```




```
linreg_small_transform <- lm((Runtime^(2/3)) ~ MatrixSize, data = df_small)
summary(linreg_small_transform)
```

```
##
## Call:
## lm(formula = (Runtime^(2/3)) ~ MatrixSize, data = df_small)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.19426 -0.14522 -0.04634  0.11432  0.27609
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9.817e-01  3.581e-02   27.42  <2e-16 ***
## MatrixSize   1.773e-02  6.852e-05   258.71  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1794 on 38 degrees of freedom
## Multiple R-squared:  0.9994, Adjusted R-squared:  0.9994
## F-statistic: 6.693e+04 on 1 and 38 DF, p-value: < 2.2e-16
```

```
#plot(linreg_small_transform)
plot(linreg_small_transform$fitted.values, linreg_small_transform$residuals)
abline(h=0, col="red")
```




```
plot(df_small$MatrixSize, df_small$Runtime^(2/3))
abline(a=linreg_small_transform$coefficients[1],
       b=linreg_small_transform$coefficients[2])
```




TPU Addition

```
df_small <- data[(data$Processor == "TPU") & (data$MatrixOperation == "Addition"),]
linreg_small <- lm(Runtime ~ MatrixSize, data = df_small)
plot(linreg_small$residuals)
abline(h=0, col="red")
```

```
boxCox(linreg_small, lambda = seq(-20, 2, 0.1))
```




```
linreg_small_transform <- lm(log10(Runtime) ~ MatrixSize, data = df_small)
summary(linreg_small_transform)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ MatrixSize, data = df_small)
##
```


```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.024742 -0.019679 -0.007998  0.002266  0.101552
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.298e+00  5.948e-03 -218.276  <2e-16 ***
## MatrixSize   1.560e-06  1.138e-05   0.137    0.892
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.02981 on 38 degrees of freedom
## Multiple R-squared:  0.0004943, Adjusted R-squared:  -0.02581
## F-statistic: 0.01879 on 1 and 38 DF, p-value: 0.8917
```

```
m <- max(abs(linreg_small_transform$residuals))
plot(linreg_small_transform$residuals, ylim=c(-m, m))
abline(h=0, col="red")
```




TPU Multiplication

```
df_small <- data[(data$Processor == "TPU") & (data$MatrixOperation == "Multiplication"),]
linreg_small <- lm(Runtime ~ MatrixSize, data = df_small)
plot(linreg_small$residuals)
abline(h=0, col="red")
```



```
boxCox(linreg_small, lambda = seq(-20, 2, 0.1))
```




```
linreg_small_transform <- lm(log10(Runtime) ~ MatrixSize, data = df_small)
summary(linreg_small_transform)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ MatrixSize, data = df_small)
##
```


```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.044167 -0.028197 -0.019721 -0.001056  0.274107
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.506e+00  1.279e-02 -117.694  <2e-16 ***
## MatrixSize   5.326e-06  2.448e-05   0.218    0.829
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06411 on 38 degrees of freedom
## Multiple R-squared:  0.001244, Adjusted R-squared: -0.02504
## F-statistic: 0.04734 on 1 and 38 DF, p-value: 0.8289
```

```
m <- max(abs(linreg_small_transform$residuals))
plot(linreg_small_transform$residuals, ylim=c(-m, m))
abline(h=0, col="red")
```




TPU Inversion

```
df_small <- data[(data$Processor == "TPU") & (data$MatrixOperation == "Inversion"),]
linreg_small <- lm(Runtime ~ MatrixSize, data = df_small)
plot(linreg_small$residuals)
abline(h=0, col="red")
```



```
boxCox(linreg_small, lambda = seq(-20, 2, 0.1))
```




```
linreg_small_transform <- lm(log10(Runtime) ~ MatrixSize, data = df_small)
summary(linreg_small_transform)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ MatrixSize, data = df_small)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.033506 -0.023723 -0.012231  0.009422  0.192046
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.705e+00  8.068e-03 -211.303  <2e-16 ***
## MatrixSize   6.430e-06  1.544e-05   0.416   0.679
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04043 on 38 degrees of freedom
## Multiple R-squared:  0.004544, Adjusted R-squared: -0.02165
## F-statistic: 0.1734 on 1 and 38 DF, p-value: 0.6794
```

```
#plot(linreg_small_transform)
m <- max(abs(linreg_small_transform$residuals))
plot(linreg_small_transform$residuals, ylim=c(-m, m))
abline(h=0, col="red")
```



ANOVA tests fixing matrix operation and matrix size

```
operation_arr <- unique(data$MatrixOperation)
size_arr <- unique(data$MatrixSize)

p_val_table <- matrix(0, nrow = length(size_arr), ncol = 2 * length(operation_arr))
rownames(p_val_table) <- size_arr
colnames(p_val_table) <- c(paste0("GPU_", operation_arr),
                           paste0("TPU_", operation_arr))
est_table <- matrix(0, nrow = length(size_arr), ncol = 2 * length(operation_arr))
ci_table <- matrix(0, nrow = length(size_arr), ncol = 2 * length(operation_arr))
```

```

for(i in 1:length(operation_arr)){
  for(j in 1:length(size_arr)){
    operation <- operation_arr[i]
    size <- size_arr[j]

    headline <- paste0("Operation = ", operation, "; Size = ", size)
    df_small <- data[(data$MatrixOperation == operation) & (data$MatrixSize == size),]
    linreg_small <- lm(Runtime ~ Processor, data = df_small)
    sm <- summary(linreg_small)
    print(paste0(headline, ": ", round(sm$adj.r.squared, 3)))
    p_val_table[j, i] <- as.numeric(sm$coefficients[2,4])
    p_val_table[j, length(operation_arr)+i] <- as.numeric(sm$coefficients[3,4])

    est_table[j, i] <- paste0(round(sm$coefficients[2,1], 3), " (",
                              round(sm$coefficients[2,2], 3), ")")
    est_table[j, length(operation_arr)+i] <- paste0(round(sm$coefficients[3,1], 3),
                                                    " (",
                                                    round(sm$coefficients[3,2], 3),
                                                    ")")


    ci_gpu <- round(sm$coefficients[2,1] +
                   c(-1,1)*qnorm(0.975)*sm$coefficients[2,2], 3)
    ci_tpu <- round(sm$coefficients[3,1] +
                   c(-1,1)*qnorm(0.975)*sm$coefficients[3,2], 3)
    ci_table[j, i] <- paste0("(", ci_gpu[1], " to ", ci_gpu[2], ")")
    ci_table[j, length(operation_arr)+i] <- paste0("(", ci_tpu[1], " to ",
                                                    ci_tpu[2], ")")

    #print(headline)
    #print(summary(linreg_small))
    m <- max(abs(linreg_small$residuals))
    plot(linreg_small$fitted.values, linreg_small$residuals, main = headline,
         xlab = "Fitted Values", ylab = "Residuals", ylim=c(-m, m))
    abline(h=0, col="red")
  }
}

```


```
## [1] "Operation = Addition; Size = 10: 0.995"
```

Operation = Addition; Size = 10




```
## [1] "Operation = Addition; Size = 20: 0.977"
```

Operation = Addition; Size = 20




```
## [1] "Operation = Addition; Size = 40: 0.993"
```


Operation = Addition; Size = 40




```
## [1] "Operation = Addition; Size = 80: 0.982"
```

Operation = Addition; Size = 80




```
## [1] "Operation = Addition; Size = 160: 0.764"
```

Operation = Addition; Size = 160




```
## [1] "Operation = Addition; Size = 320: 0.999"
```

Operation = Addition; Size = 320




```
## [1] "Operation = Addition; Size = 640: 1"
```

Operation = Addition; Size = 640




```
## [1] "Operation = Addition; Size = 1280: 1"
```

Operation = Addition; Size = 1280




```
## [1] "Operation = Multiplication; Size = 10: 0.904"
```

Operation = Multiplication; Size = 10




```
## [1] "Operation = Multiplication; Size = 20: 0.873"
```

Operation = Multiplication; Size = 20




```
## [1] "Operation = Multiplication; Size = 40: 0.991"
```

Operation = Multiplication; Size = 40




```
## [1] "Operation = Multiplication; Size = 80: 0.991"
```

Operation = Multiplication; Size = 80




```
## [1] "Operation = Multiplication; Size = 160: 1"
```

Operation = Multiplication; Size = 160




```
## [1] "Operation = Multiplication; Size = 320: 1"
```

Operation = Multiplication; Size = 320




```
## [1] "Operation = Multiplication; Size = 640: 1"
```

Operation = Multiplication; Size = 640




```
## [1] "Operation = Multiplication; Size = 1280: 1"
```

Operation = Multiplication; Size = 1280




```
## [1] "Operation = Inversion; Size = 10: 0.999"
```

Operation = Inversion; Size = 10




```
## [1] "Operation = Inversion; Size = 20: 1"
```

Operation = Inversion; Size = 20




```
## [1] "Operation = Inversion; Size = 40: 1"
```


Operation = Inversion; Size = 40




```
## [1] "Operation = Inversion; Size = 80: 1"
```

Operation = Inversion; Size = 80




```
## [1] "Operation = Inversion; Size = 160: 1"
```

Operation = Inversion; Size = 160




```
## [1] "Operation = Inversion; Size = 320: 1"
```

Operation = Inversion; Size = 320




```
## [1] "Operation = Inversion; Size = 640: 1"
```

Operation = Inversion; Size = 640



```
## [1] "Operation = Inversion; Size = 1280: 1"
```

Operation = Inversion; Size = 1280



```

#print(p_val_table)

parse_number <- function(x){
  ret <- paste0("$", str_replace_all(formatC(x, digits=4), "e", " \\times 10^{")
  if(str_detect(ret, "\\times")){
    ret <- paste0(ret, "}")
  }
  ret <- paste0(ret, "$")
  ret
}

p_val_display <- p_val_table
for(i in 1:nrow(p_val_display)){
  for(j in 1:ncol(p_val_display)){
    p_val_display[i, j] <- parse_number(p_val_table[i, j])
  }
}
#p_val_display %>% print()

for(i in 1:nrow(p_val_display)){
  #paste0(p_val_display[i,], collapse = " & ") %>% print()
  #paste0(ci_table[i,], collapse = " & ") %>% print()
}

```

```

df_small <- data[(data$MatrixOperation == "Inversion") & (data$MatrixSize == 1280),]
linreg_small <- lm(Runtime ~ Processor, data = df_small)
summary(linreg_small)

```

```

##
## Call:
## lm(formula = Runtime ~ Processor, data = df_small)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.7751 -0.0557 -0.0008  0.0129  5.7155
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   908.077      1.289   704.7  <2e-16 ***
## ProcessorGPU -794.204      1.822  -435.8  <2e-16 ***
## ProcessorTPU -908.058      1.822  -498.3  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.881 on 12 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 1.474e+05 on 2 and 12 DF, p-value: < 2.2e-16


```

Big model


```

linreg_big <- lm(Runtime ~ MatrixOperation + Processor + MatrixSize, data = data)
plot(linreg_big$residuals)
abline(h=0, col="red")

```



```
boxCox(linreg_big)
```



```
linreg_big_trans <- lm(Runtime^(-1/3) ~ MatrixOperation * Processor +
  poly(MatrixSize, 4),
  data = data)
summary(linreg_big_trans)
```


```
##
## Call:
```

```

## lm(formula = Runtime^(-1/3) ~ MatrixOperation * Processor + poly(MatrixSize,
##    4), data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.14670 -0.33016 -0.06913  0.38243  1.17423
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   2.30016    0.07962  28.888 < 2e-16
## MatrixOperationInversion      -1.36533    0.11260 -12.125 < 2e-16
## MatrixOperationMultiplication -0.82850    0.11260  -7.358 1.37e-12
## ProcessorGPU                  -0.04793    0.11260  -0.426 0.670607
## ProcessorTPU                   0.40803    0.11260   3.624 0.000334
## poly(MatrixSize, 4)1          -7.84492    0.50358 -15.578 < 2e-16
## poly(MatrixSize, 4)2           3.46485    0.50358   6.880 2.80e-11
## poly(MatrixSize, 4)3          -1.73253    0.50358  -3.440 0.000652
## poly(MatrixSize, 4)4           0.33158    0.50358   0.658 0.510688
## MatrixOperationInversion:ProcessorGPU -0.30152    0.15925  -1.893 0.059131
## MatrixOperationMultiplication:ProcessorGPU 0.35805    0.15925   2.248 0.025177
## MatrixOperationInversion:ProcessorTPU    2.35357    0.15925  14.779 < 2e-16
## MatrixOperationMultiplication:ProcessorTPU 1.29535    0.15925   8.134 7.46e-15
##
## (Intercept)                ***
## MatrixOperationInversion    ***
## MatrixOperationMultiplication ***
## ProcessorGPU
## ProcessorTPU                ***
## poly(MatrixSize, 4)1        ***
## poly(MatrixSize, 4)2        ***
## poly(MatrixSize, 4)3        ***
## poly(MatrixSize, 4)4
## MatrixOperationInversion:ProcessorGPU .
## MatrixOperationMultiplication:ProcessorGPU *
## MatrixOperationInversion:ProcessorTPU ***
## MatrixOperationMultiplication:ProcessorTPU ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5036 on 347 degrees of freedom
## Multiple R-squared:  0.8226, Adjusted R-squared:  0.8165
## F-statistic: 134.1 on 12 and 347 DF, p-value: < 2.2e-16

plot(linreg_big_trans$residuals)
abline(h=0, col="red")

```



Lixian Chen Part I

When the processor is the same, is there any operation and matrix size effect? We want to answer the following question: in each scenario, which processor should we use? When the matrix size is the same, is there any processor or operation type effect, or interaction?

Getting Data Ready

```
# Read Data
df <- fread("../data/Runtime.csv", header=TRUE)
attach(df)

MatrixOperation<-factor(df$MatrixOperation)
Processor<-factor(df$Processor)

summary(Runtime)
```

```
##      Min.   1st Qu.   Median     Mean  3rd Qu.    Max.
##  0.0183   0.0313   0.0712  28.2915   1.2407  913.7929
```

```
summary(Processor)
```

```
## CPU GPU TPU
## 120 120 120
```


```
summary(MatrixOperation)
```

```
##      Addition      Inversion Multiplication
##      120          120          120
```


```
summary(MatrixSize)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      10.0   35.0   120.0   318.8   400.0  1280.0
```

```
op=par(mfrow=c(1,1))
#op=par(mfrow=c(2,1))
interaction.plot(df$MatrixSize, df$MatrixOperation, df$Runtime, ylab = "Runtime (seconds)",
                 xlab = "Matrix Size", trace.label = "Matrix Operation", col = c("blue4", "red4", "green4"))
```



```
interaction.plot(df$MatrixSize, df$MatrixOperation, log(df$Runtime), ylab = "Runtime (seconds)",
                 xlab = "Matrix Size", trace.label = "Matrix Operation", col = c("blue4", "red4", "green4"))
```

```
par(op)
```

```
CPUdata <- df %>% dplyr::filter(Processor=="CPU")
GPUdata <- df %>% dplyr::filter(Processor=="GPU")
TPUdata <- df %>% dplyr::filter(Processor=="TPU")
```

```
summary(CPUdata$Runtime)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##  0.0185   0.0576   0.5576   76.2816  13.0731  913.7929
```

```
summary(GPUdata$Runtime)
```

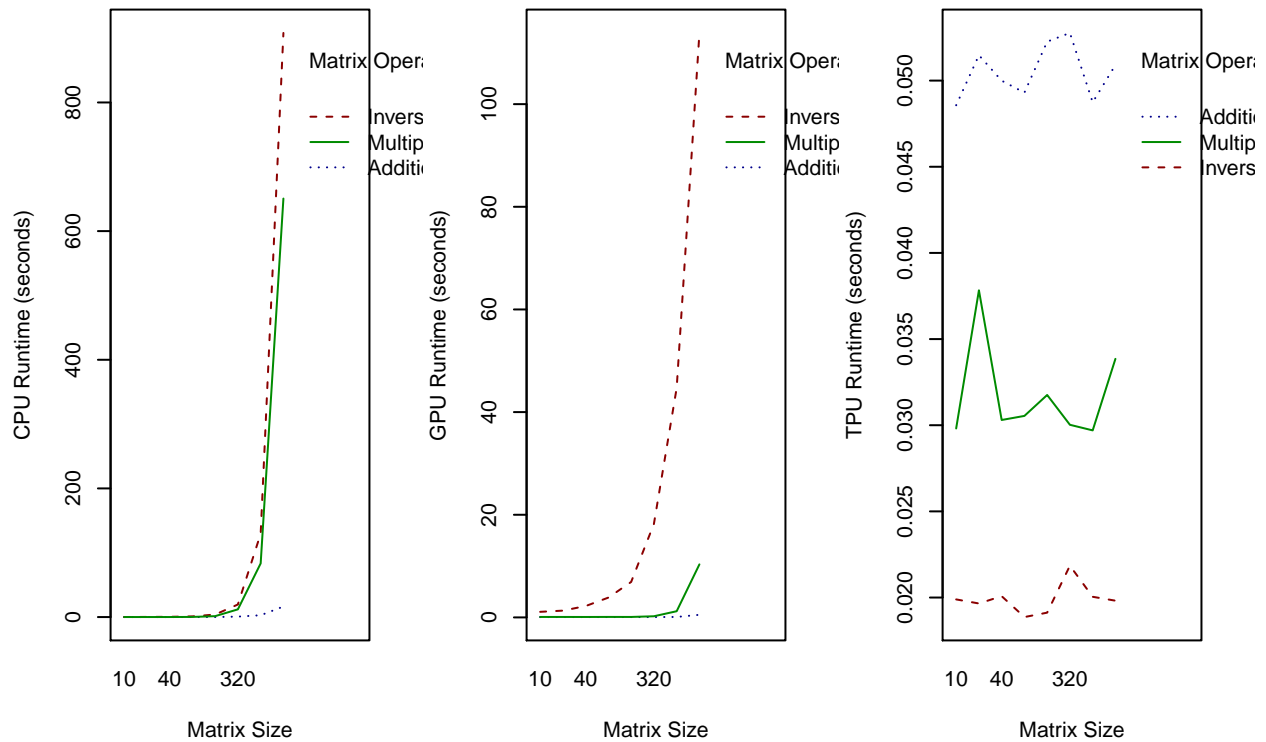
```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##  0.06736   0.08084   0.15957   8.55880   2.63984  113.99131
```

```
summary(TPUdata$Runtime)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##  0.01831   0.02046   0.03004   0.03405   0.04831   0.06362
```

```
op=par(mfrow=c(1,3))
```


```
interaction.plot(CPUdata$MatrixSize, CPUdata$MatrixOperation, CPUdata$Runtime, ylab = "CPU Runtime (seconds)",
                  xlab = "Matrix Size", trace.label = "Matrix Operation", col = c("blue4", "red4", "green4"))
interaction.plot(GPUdata$MatrixSize, GPUdata$MatrixOperation, GPUdata$Runtime, ylab = "GPU Runtime (seconds)",
                  xlab = "Matrix Size", trace.label = "Matrix Operation", col = c("blue4", "red4", "green4"))
interaction.plot(TPUdata$MatrixSize, TPUdata$MatrixOperation, TPUdata$Runtime, ylab = "TPU Runtime (seconds)",
                  xlab = "Matrix Size", trace.label = "Matrix Operation", col = c("blue4", "red4", "green4"))
```



```
par(op)

#boxplot(Runtime~Processor*MatrixOperation)
#tapply(Runtime,list(Processor, MatrixOperation),mean)
#tapply(Runtime,MatrixOperation,mean)

boxplot(Runtime~Processor*MatrixOperation)
```



```
round(tapply(Runtime,list(Processor, MatrixOperation),mean),digits=4)
```

```
##      Addition Inversion Multiplication
## CPU   2.5147  132.8553      93.4747
## GPU   0.1280   24.0220      1.5264
## TPU   0.0505   0.0199      0.0317
```

```
#tapply(Runtime,MatrixOperation,mean)
```

```
# group_by(df, Processor) %>%
#   summarise(
#     count = n(),
#     mean = mean(Runtime, na.rm = TRUE),
#     sd = sd(Runtime, na.rm = TRUE)
#   )
#
# group_by(df, MatrixOperation) %>%
#   summarise(
#     count = n(),
#     mean = mean(Runtime, na.rm = TRUE),
#     sd = sd(Runtime, na.rm = TRUE)
#   )
```

```
detach(df)
```

```
size320data <- df %>% dplyr::filter(MatrixSize==320)
size640data <- df %>% dplyr::filter(MatrixSize==640)
size1280data <- df %>% dplyr::filter(MatrixSize==1280)
```

Mean of Runtime and Boxplot

```
tapply(size320data$Runtime,list(size320data$Processor, size320data$MatrixOperation),mean)
```

```
##      Addition Inversion Multiplication
## CPU 0.71350303 19.54677558  11.97371540
## GPU 0.07051702 18.21737680   0.22623463
## TPU 0.05276523  0.02184839   0.03002601
```

```
tapply(size640data$Runtime,list(size640data$Processor, size640data$MatrixOperation),mean)
```


```
##      Addition Inversion Multiplication
## CPU 2.84506497 129.50422368  83.41594014
## GPU 0.08734841  44.69854727   1.20817008
## TPU 0.04873762  0.02004447   0.02969995
```

```
tapply(size1280data$Runtime,list(size1280data$Processor, size1280data$MatrixOperation),mean)
```

```
##      Addition Inversion Multiplication
## CPU 16.39168239 908.07733464 650.62470641
## GPU 0.51256037 113.87343698  10.33359056
## TPU 0.05090575  0.01981788   0.03385234
```


```
boxplot(Runtime~Processor*MatrixOperation, data = size320data, main="At the level of matrix size=320")
```

At the level of matrix size=320




```
boxplot(Runtime~Processor*MatrixOperation, data = size640data, main="At the level of matrix size=640")
```

At the level of matrix size=640



```
boxplot(Runtime~Processor*MatrixOperation, data = size1280data, main="At the level of matrix size=1280")
```



Analysis

Anova

```
##### 320
fit2<-lm(log10(Runtime)~Processor+MatrixOperation, data = size320data)
summary(fit2)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.84717 -0.41443 -0.05434  0.39193  0.83638
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.1912    0.1870   1.022  0.3128
## ProcessorGPU     -0.9198    0.2048  -4.491 5.90e-05 ***
## ProcessorTPU     -2.2309    0.2048 -10.892 1.55e-13 ***
## MatrixOperationInversion  1.1534    0.2048   5.631 1.56e-06 ***
## MatrixOperationMultiplication  0.4957    0.2048   2.420  0.0201 *
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5609 on 40 degrees of freedom
## Multiple R-squared:  0.7914, Adjusted R-squared:  0.7706
## F-statistic: 37.94 on 4 and 40 DF,  p-value: 4.091e-13

fit1<-lm(log10(Runtime)~Processor*MatrixOperation, data = size320data)
summary(fit1)

##
## Call:
## lm(formula = log10(Runtime) ~ Processor * MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.064926 -0.008334 -0.000891  0.002071  0.157877
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   -0.14669     0.01471  -9.971 6.71e-12
## ProcessorGPU                   -1.00513     0.02081 -48.310 < 2e-16
## ProcessorTPU                   -1.13204     0.02081 -54.409 < 2e-16
## MatrixOperationInversion        1.43776     0.02081  69.104 < 2e-16
## MatrixOperationMultiplication    1.22492     0.02081  58.874 < 2e-16
## ProcessorGPU:MatrixOperationInversion  0.97455     0.02942  33.121 < 2e-16
## ProcessorTPU:MatrixOperationInversion -1.82762     0.02942 -62.113 < 2e-16
## ProcessorGPU:MatrixOperationMultiplication -0.71857     0.02942 -24.421 < 2e-16
## ProcessorTPU:MatrixOperationMultiplication -1.46896     0.02942 -49.924 < 2e-16
##
## (Intercept)                    ***
## ProcessorGPU                    ***
## ProcessorTPU                    ***
## MatrixOperationInversion        ***
## MatrixOperationMultiplication    ***
## ProcessorGPU:MatrixOperationInversion ***
## ProcessorTPU:MatrixOperationInversion ***
## ProcessorGPU:MatrixOperationMultiplication ***
## ProcessorTPU:MatrixOperationMultiplication ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0329 on 36 degrees of freedom
## Multiple R-squared:  0.9994, Adjusted R-squared:  0.9992
## F-statistic: 6965 on 8 and 36 DF,  p-value: < 2.2e-16

anova(fit2, fit1)

## Analysis of Variance Table
##
## Model 1: log10(Runtime) ~ Processor + MatrixOperation
## Model 2: log10(Runtime) ~ Processor * MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      40 12.586
```

```
## 2      36 0.039 4      12.547 2898.4 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

mod320<-aov(log10(Runtime)~Processor+MatrixOperation, data = size320data)
#anova(mod320)
Anova(mod320,type="III")

## Anova Table (Type III tests)
##
## Response: log10(Runtime)
##              Sum Sq Df  F value    Pr(>F)
## (Intercept)      0.1076  1    99.421 6.713e-12 ***
## Processor        3.8465  2  1777.140 < 2.2e-16 ***
## MatrixOperation   6.0215  2  2782.019 < 2.2e-16 ***
## Processor:MatrixOperation 12.5469  4 2898.441 < 2.2e-16 ***
## Residuals         0.0390 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

modF320<-lm(log10(Runtime)~Processor+MatrixOperation, data = size320data)
modA320<-lm(log10(Runtime)~Processor, data = size320data)
modB320<-lm(log10(Runtime)~MatrixOperation, data = size320data)

anova(modA320, modF320)

## Analysis of Variance Table
##
## Model 1: log10(Runtime) ~ Processor
## Model 2: log10(Runtime) ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 22.629
## 2      40 12.586  2    10.043 15.959 8.024e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

anova(modB320, modF320)

## Analysis of Variance Table
##
## Model 1: log10(Runtime) ~ MatrixOperation
## Model 2: log10(Runtime) ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 50.295
## 2      40 12.586  2    37.71 59.923 9.271e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

##### 640
fit2_640<-lm(Runtime~Processor+MatrixOperation, data = size640data)
summary(fit2)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.84717 -0.41443 -0.05434  0.39193  0.83638
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.1912     0.1870   1.022  0.3128
## ProcessorGPU     -0.9198     0.2048  -4.491 5.90e-05 ***
## ProcessorTPU     -2.2309     0.2048 -10.892 1.55e-13 ***
## MatrixOperationInversion  1.1534     0.2048   5.631 1.56e-06 ***
## MatrixOperationMultiplication  0.4957     0.2048   2.420  0.0201 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5609 on 40 degrees of freedom
## Multiple R-squared:  0.7914, Adjusted R-squared:  0.7706
## F-statistic: 37.94 on 4 and 40 DF,  p-value: 4.091e-13

fit1_640<-lm(Runtime~Processor*MatrixOperation, data = size640data)
summary(fit1_640)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation, data = size640data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04780 -0.00908 -0.00014  0.00088  2.65489
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)      2.8451     0.2781  10.229
## ProcessorGPU     -2.7577     0.3933   -7.011
## ProcessorTPU     -2.7963     0.3933   -7.109
## MatrixOperationInversion 126.6592     0.3933 322.008
## MatrixOperationMultiplication  80.5709     0.3933 204.837
## ProcessorGPU:MatrixOperationInversion -82.0480     0.5563 -147.497
## ProcessorTPU:MatrixOperationInversion -126.6879     0.5563 -227.746
## ProcessorGPU:MatrixOperationMultiplication -79.4501     0.5563 -142.827
## ProcessorTPU:MatrixOperationMultiplication -80.5899     0.5563 -144.876
##              Pr(>|t|)
## (Intercept)      3.38e-12 ***
## ProcessorGPU      3.18e-08 ***
## ProcessorTPU      2.36e-08 ***
## MatrixOperationInversion < 2e-16 ***
## MatrixOperationMultiplication < 2e-16 ***
## ProcessorGPU:MatrixOperationInversion < 2e-16 ***
## ProcessorTPU:MatrixOperationInversion < 2e-16 ***
## ProcessorGPU:MatrixOperationMultiplication < 2e-16 ***
## ProcessorTPU:MatrixOperationMultiplication < 2e-16 ***
```



```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6219 on 36 degrees of freedom
## Multiple R-squared:  0.9998, Adjusted R-squared:  0.9998
## F-statistic: 2.928e+04 on 8 and 36 DF,  p-value: < 2.2e-16
```

```
anova(fit2_640, fit1_640)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor + MatrixOperation
## Model 2: Runtime ~ Processor * MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      40 23128.8
## 2      36   13.9  4      23115 14940 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod640<-aov(Runtime~Processor*MatrixOperation, data = size640data)
Anova(mod640,type="III")
```

```
## Anova Table (Type III tests)
##
## Response: Runtime
##
##              Sum Sq Df    F value    Pr(>F)
## (Intercept)      40  1    104.635 3.380e-12 ***
## Processor        26  2     33.235 6.649e-09 ***
## MatrixOperation 41097  2 53125.511 < 2.2e-16 ***
## Processor:MatrixOperation 23115  4 14940.049 < 2.2e-16 ***
## Residuals        14 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
modF640<-lm(Runtime~Processor+MatrixOperation, data = size640data)
modA640<-lm(Runtime~Processor, data = size640data)
modB640<-lm(Runtime~MatrixOperation, data = size640data)

anova(modA640, modF640)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 47583
## 2      40 23129  2      24454 21.146 5.421e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(modB640, modF640)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ MatrixOperation
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df  RSS Df Sum of Sq    F    Pr(>F)
## 1      42 66151
## 2      40 23129  2      43023 37.203 7.452e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##### 1280
```

```
fit2_1280<-lm(Runtime~Processor+MatrixOperation, data = size1280data)
summary(fit2)
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ Processor + MatrixOperation, data = size320data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.84717 -0.41443 -0.05434  0.39193  0.83638
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.1912     0.1870   1.022   0.3128
## ProcessorGPU     -0.9198     0.2048  -4.491 5.90e-05 ***
## ProcessorTPU     -2.2309     0.2048 -10.892 1.55e-13 ***
## MatrixOperationInversion  1.1534     0.2048   5.631 1.56e-06 ***
## MatrixOperationMultiplication  0.4957     0.2048   2.420  0.0201 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5609 on 40 degrees of freedom
## Multiple R-squared:  0.7914, Adjusted R-squared:  0.7706
## F-statistic: 37.94 on 4 and 40 DF,  p-value: 4.091e-13
```

```
fit1_1280<-lm(Runtime~Processor*MatrixOperation, data = size1280data)
summary(fit1_1280)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation, data = size1280data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.7751 -0.0345 -0.0008  0.0132  5.7155
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)      16.3917     0.8192  20.01
```

```
## ProcessorGPU -15.8791 1.1585 -13.71
## ProcessorTPU -16.3408 1.1585 -14.11
## MatrixOperationInversion 891.6857 1.1585 769.69
## MatrixOperationMultiplication 634.2330 1.1585 547.46
## ProcessorGPU:MatrixOperationInversion -778.3248 1.6384 -475.06
## ProcessorTPU:MatrixOperationInversion -891.7167 1.6384 -544.27
## ProcessorGPU:MatrixOperationMultiplication -624.4120 1.6384 -381.12
## ProcessorTPU:MatrixOperationMultiplication -634.2501 1.6384 -387.12
## Pr(>|t|)
## (Intercept) < 2e-16 ***
## ProcessorGPU 7.42e-16 ***
## ProcessorTPU 3.09e-16 ***
## MatrixOperationInversion < 2e-16 ***
## MatrixOperationMultiplication < 2e-16 ***
## ProcessorGPU:MatrixOperationInversion < 2e-16 ***
## ProcessorTPU:MatrixOperationInversion < 2e-16 ***
## ProcessorGPU:MatrixOperationMultiplication < 2e-16 ***
## ProcessorTPU:MatrixOperationMultiplication < 2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.832 on 36 degrees of freedom
## Multiple R-squared: 1, Adjusted R-squared: 1
## F-statistic: 1.751e+05 on 8 and 36 DF, p-value: < 2.2e-16
```

```
anova(fit2_1280, fit1_1280)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor + MatrixOperation
## Model 2: Runtime ~ Processor * MatrixOperation
## Res.Df RSS Df Sum of Sq F Pr(>F)
## 1 40 1281658
## 2 36 121 4 1281538 95485 < 2.2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
modF1280<-lm(Runtime~Processor+MatrixOperation, data = size1280data)
modA1280<-lm(Runtime~Processor, data = size1280data)
modB1280<-lm(Runtime~MatrixOperation, data = size1280data)
anova(modA1280, modF1280)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ Processor
## Model 2: Runtime ~ Processor + MatrixOperation
## Res.Df RSS Df Sum of Sq F Pr(>F)
## 1 42 2145628
## 2 40 1281658 2 863970 13.482 3.345e-05 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(modB1280, modF1280)
```

```
## Analysis of Variance Table
##
## Model 1: Runtime ~ MatrixOperation
## Model 2: Runtime ~ Processor + MatrixOperation
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      42 3837050
## 2      40 1281658  2   2555392 39.876 2.989e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod1280<-aov(Runtime~Processor*MatrixOperation, data = size1280data)
Anova(mod1280,type="III")
```

```
## Anova Table (Type III tests)
##
## Response: Runtime
##
##               Sum Sq Df    F value    Pr(>F)
## (Intercept)      1343  1    400.39 < 2.2e-16 ***
## Processor         866  2    128.99 < 2.2e-16 ***
## MatrixOperation 2106061  2 313838.85 < 2.2e-16 ***
## Processor:MatrixOperation 1281538  4  95485.43 < 2.2e-16 ***
## Residuals         121 36
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#also fix processor, 1 way ANOVA
size320CPUdata <- data %>% dplyr::filter(MatrixSize==320 & Processor=="CPU")
size640CPUdata <- data %>% dplyr::filter(MatrixSize==640 & Processor=="CPU")
size1280CPUdata <- data %>% dplyr::filter(MatrixSize==1280 & Processor=="CPU")

size320GPUdata <- data %>% dplyr::filter(MatrixSize==320 & Processor=="GPU")
size640GPUdata <- data %>% dplyr::filter(MatrixSize==640 & Processor=="GPU")
size1280GPUdata <- data %>% filter(MatrixSize==1280 & Processor=="GPU")

size320TPUdata <- data %>% dplyr::filter(MatrixSize==320 & Processor=="TPU")
size640TPUdata <- data %>% dplyr::filter(MatrixSize==640 & Processor=="TPU")
size1280TPUdata <- data %>% dplyr::filter(MatrixSize==1280 & Processor=="TPU")

mod320CPU<-lm(Runtime~MatrixOperation, data = size320CPUdata)
summary(mod320CPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size320CPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.135726 -0.025213 -0.004924  0.022535  0.166785
##
## Coefficients:
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.71350    0.03425   20.83 8.67e-11 ***
## MatrixOperationInversion  18.83327    0.04844  388.78 < 2e-16 ***
## MatrixOperationMultiplication 11.26021    0.04844  232.45 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.07659 on 12 degrees of freedom
## Multiple R-squared:  0.9999, Adjusted R-squared:  0.9999
## F-statistic: 7.654e+04 on 2 and 12 DF, p-value: < 2.2e-16
```

```
mod320GPU<-lm(Runtime~MatrixOperation, data = size320GPUdata)
summary(mod320GPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size320GPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.037341 -0.001891 -0.000090  0.003157  0.027248
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.070517   0.006911   10.20 2.88e-07 ***
## MatrixOperationInversion  18.146860   0.009774 1856.69 < 2e-16 ***
## MatrixOperationMultiplication 0.155718   0.009774   15.93 1.95e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01545 on 12 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 2.279e+06 on 2 and 12 DF, p-value: < 2.2e-16
```

```
mod320TPU<-lm(Runtime~MatrixOperation, data = size320TPUdata)
summary(mod320TPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size320TPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.004670 -0.001985 -0.001053  0.001259  0.009004
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.052765   0.001724  30.613 9.26e-13 ***
## MatrixOperationInversion  -0.030917   0.002438 -12.684 2.60e-08 ***
## MatrixOperationMultiplication -0.022739   0.002438  -9.329 7.55e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 0.003854 on 12 degrees of freedom
## Multiple R-squared:  0.9351, Adjusted R-squared:  0.9242
## F-statistic: 86.39 on 2 and 12 DF,  p-value: 7.504e-08
```

```
mod640CPU<-lm(Runtime~MatrixOperation, data = size640CPUdata)
summary(mod640CPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size640CPUdata)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-1.04780	-0.67767	-0.04855	0.03665	2.65489

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.8451	0.4817	5.906	7.18e-05 ***
MatrixOperationInversion	126.6592	0.6812	185.922	< 2e-16 ***
MatrixOperationMultiplication	80.5709	0.6812	118.269	< 2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.077 on 12 degrees of freedom
## Multiple R-squared:  0.9997, Adjusted R-squared:  0.9996
## F-statistic: 1.771e+04 on 2 and 12 DF,  p-value: < 2.2e-16
```

```
mod640GPU<-lm(Runtime~MatrixOperation, data = size640GPUdata)
summary(mod640GPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size640GPUdata)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.0296608	-0.0026586	-0.0000269	0.0062409	0.0174795

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.087348	0.005132	17.02	9.09e-10 ***
MatrixOperationInversion	44.611199	0.007258	6146.63	< 2e-16 ***
MatrixOperationMultiplication	1.120822	0.007258	154.43	< 2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01148 on 12 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 2.457e+07 on 2 and 12 DF,  p-value: < 2.2e-16
```

```
mod640TPU<-lm(Runtime~MatrixOperation, data = size640TPUdata)
summary(mod640TPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size640TPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0013937 -0.0007722  0.0003468  0.0007116  0.0008858
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.0487376   0.0003826   127.39 < 2e-16 ***
## MatrixOperationInversion -0.0286932   0.0005411  -53.03 1.33e-15 ***
## MatrixOperationMultiplication -0.0190377   0.0005411  -35.19 1.77e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0008555 on 12 degrees of freedom
## Multiple R-squared:  0.9959, Adjusted R-squared:  0.9952
## F-statistic: 1456 on 2 and 12 DF, p-value: 4.774e-15
```

```
mod1280CPU<-lm(Runtime~MatrixOperation, data = size1280CPUdata)
summary(mod1280CPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size1280CPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.7751 -2.1233 -0.0388  1.3311  5.7155
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      16.392      1.416   11.58 7.19e-08 ***
## MatrixOperationInversion  891.686      2.002  445.43 < 2e-16 ***
## MatrixOperationMultiplication 634.233      2.002  316.82 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.165 on 12 degrees of freedom
## Multiple R-squared:  0.9999, Adjusted R-squared:  0.9999
## F-statistic: 1.051e+05 on 2 and 12 DF, p-value: < 2.2e-16
```

```
mod1280GPU<-lm(Runtime~MatrixOperation, data = size1280GPUdata)
summary(mod1280GPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size1280GPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.44634 -0.03290 -0.00015  0.01275  0.40972
```

```
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.51256    0.09736   5.265   2e-04 ***
## MatrixOperationInversion 113.36088    0.13768 823.339 <2e-16 ***
## MatrixOperationMultiplication 9.82103    0.13768  71.330 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2177 on 12 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 4.162e+05 on 2 and 12 DF, p-value: < 2.2e-16
```

```
mod1280TPU<-lm(Runtime~MatrixOperation, data = size1280TPUdata)
summary(mod1280TPU)
```

```
##
## Call:
## lm(formula = Runtime ~ MatrixOperation, data = size1280TPUdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0052021 -0.0029268 -0.0008882  0.0003531  0.0168402
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.050906    0.002731  18.640 3.17e-10 ***
## MatrixOperationInversion -0.031088    0.003862  -8.049 3.53e-06 ***
## MatrixOperationMultiplication -0.017053    0.003862  -4.415 0.000842 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.006107 on 12 degrees of freedom
## Multiple R-squared:  0.8441, Adjusted R-squared:  0.8182
## F-statistic: 32.5 on 2 and 12 DF, p-value: 1.434e-05
```

```
#grand overall model
unique(data$MatrixSize)
```

```
## [1]  10  20  40  80 160 320 640 1280
```

```
mod<-aov(Runtime~Processor+MatrixOperation+as.factor(MatrixSize), data = df)
summary(mod)
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## Processor      2  418909   209455   17.292 6.92e-08 ***
## MatrixOperation  2   160590    80295    6.629  0.00149 **
## as.factor(MatrixSize) 7  1356897   193842   16.003 < 2e-16 ***
## Residuals     348  4215313    12113
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



```
mod2<-lm(Runtime~Processor+MatrixOperation+as.factor(MatrixSize), data = df)
summary(mod2)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor + MatrixOperation + as.factor(MatrixSize),
##     data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -193.22  -48.01    2.87   24.60  652.92
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    20.76253   20.09392   1.033 0.302193
## ProcessorGPU   -67.72278   14.20855  -4.766 2.76e-06 ***
## ProcessorTPU   -76.24754   14.20855  -5.366 1.47e-07 ***
## MatrixOperationInversion  51.40131   14.20855   3.618 0.000341 ***
## MatrixOperationMultiplication 30.77990   14.20855   2.166 0.030967 *
## as.factor(MatrixSize)20     0.04295   23.20246   0.002 0.998524
## as.factor(MatrixSize)40     0.16240   23.20246   0.007 0.994420
## as.factor(MatrixSize)80     0.43837   23.20246   0.019 0.984937
## as.factor(MatrixSize)160    1.23193   23.20246   0.053 0.957687
## as.factor(MatrixSize)320    5.48415   23.20246   0.236 0.813292
## as.factor(MatrixSize)640   28.92915   23.20246   1.247 0.213304
## as.factor(MatrixSize)1280 188.71361   23.20246   8.133 7.45e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 110.1 on 348 degrees of freedom
## Multiple R-squared:  0.3148, Adjusted R-squared:  0.2931
## F-statistic: 14.53 on 11 and 348 DF,  p-value: < 2.2e-16
```

```
#res <- resid(mod2)
#plot(fitted(mod2), res)
#abline(0,0)
mod3<-lm(Runtime~Processor*MatrixOperation+as.factor(MatrixSize), data = df)
summary(mod3)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation + as.factor(MatrixSize),
##     data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -160.59  -60.48   21.96   27.09  620.35
##
## Coefficients:
##              Estimate Std. Error t value
## (Intercept)   -25.61060   22.74453  -1.126
## ProcessorGPU   -2.38672   24.12422  -0.099
```

```
## ProcessorTPU -2.46422 24.12422 -0.102
## MatrixOperationInversion 130.34056 24.12422 5.403
## MatrixOperationMultiplication 90.96003 24.12422 3.770
## as.factor(MatrixSize)20 0.04295 22.74453 0.002
## as.factor(MatrixSize)40 0.16240 22.74453 0.007
## as.factor(MatrixSize)80 0.43837 22.74453 0.019
## as.factor(MatrixSize)160 1.23193 22.74453 0.054
## as.factor(MatrixSize)320 5.48415 22.74453 0.241
## as.factor(MatrixSize)640 28.92915 22.74453 1.272
## as.factor(MatrixSize)1280 188.71361 22.74453 8.297
## ProcessorGPU:MatrixOperationInversion -106.44661 34.11680 -3.120
## ProcessorTPU:MatrixOperationInversion -130.37115 34.11680 -3.821
## ProcessorGPU:MatrixOperationMultiplication -89.56157 34.11680 -2.625
## ProcessorTPU:MatrixOperationMultiplication -90.97880 34.11680 -2.667
## Pr(>|t|)
## (Intercept) 0.260946
## ProcessorGPU 0.921248
## ProcessorTPU 0.918699
## MatrixOperationInversion 1.23e-07 ***
## MatrixOperationMultiplication 0.000192 ***
## as.factor(MatrixSize)20 0.998495
## as.factor(MatrixSize)40 0.994307
## as.factor(MatrixSize)80 0.984634
## as.factor(MatrixSize)160 0.956836
## as.factor(MatrixSize)320 0.809606
## as.factor(MatrixSize)640 0.204262
## as.factor(MatrixSize)1280 2.45e-15 ***
## ProcessorGPU:MatrixOperationInversion 0.001961 **
## ProcessorTPU:MatrixOperationInversion 0.000158 ***
## ProcessorGPU:MatrixOperationMultiplication 0.009048 **
## ProcessorTPU:MatrixOperationMultiplication 0.008022 **
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 107.9 on 344 degrees of freedom
## Multiple R-squared: 0.3491, Adjusted R-squared: 0.3207
## F-statistic: 12.3 on 15 and 344 DF, p-value: < 2.2e-16
```

```
modN<-aov(Runtime~Processor+MatrixOperation+MatrixSize, data = df)
summary(modN)
```

```
##          Df Sum Sq Mean Sq F value    Pr(>F)
## Processor    2  418909   209455   16.915 9.64e-08 ***
## MatrixOperation  2  160590    80295    6.484 0.00172 **
## MatrixSize     1 1188587 1188587   95.984 < 2e-16 ***
## Residuals   354 4383623   12383
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
mod2N<-lm(Runtime~Processor+MatrixOperation+MatrixSize, data = df)
summary(mod2N)
```

```
##
```

```
## Call:
## lm(formula = Runtime ~ Processor + MatrixOperation + MatrixSize,
##     data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -166.02  -49.26   -6.72   27.83  680.12
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      4.65633    13.86977   0.336 0.737283
## ProcessorGPU     -67.72278    14.36612  -4.714 3.50e-06 ***
## ProcessorTPU     -76.24754    14.36612  -5.307 1.97e-07 ***
## MatrixOperationInversion  51.40131    14.36612   3.578 0.000395 ***
## MatrixOperationMultiplication 30.77990    14.36612   2.143 0.032832 *
## MatrixSize        0.13877     0.01416   9.797 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 111.3 on 354 degrees of freedom
## Multiple R-squared:  0.2874, Adjusted R-squared:  0.2773
## F-statistic: 28.56 on 5 and 354 DF,  p-value: < 2.2e-16
```

```
#res <- resid(mod2N)
#plot(fitted(mod2N), res)
#abline(0,0)
```

```
modNI<-aov(Runtime~Processor*MatrixOperation+MatrixSize, data = df)
summary(modNI)
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## Processor      2  418909   209455   17.570 5.35e-08 ***
## MatrixOperation 2  160590    80295    6.736 0.00135 **
## MatrixSize      1 1188587  1188587   99.706 < 2e-16 ***
## Processor:MatrixOperation 4  211306    52826    4.431 0.00166 **
## Residuals     350 4172318    11921
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
mod2NI<-lm(Runtime~Processor*MatrixOperation+MatrixSize, data = df)
summary(mod2NI)
```

```
##
## Call:
## lm(formula = Runtime ~ Processor * MatrixOperation + MatrixSize,
##     data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -133.39  -50.94    8.95   37.59  647.55
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      4.65633    13.86977   0.336 0.737283
## ProcessorGPU     -67.72278    14.36612  -4.714 3.50e-06 ***
## ProcessorTPU     -76.24754    14.36612  -5.307 1.97e-07 ***
## MatrixOperationInversion  51.40131    14.36612   3.578 0.000395 ***
## MatrixOperationMultiplication 30.77990    14.36612   2.143 0.032832 *
## MatrixSize        0.13877     0.01416   9.797 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
## (Intercept) -41.7168 17.8226 -2.341
## ProcessorGPU -2.3867 24.4140 -0.098
## ProcessorTPU -2.4642 24.4140 -0.101
## MatrixOperationInversion 130.3406 24.4140 5.339
## MatrixOperationMultiplication 90.9600 24.4140 3.726
## MatrixSize 0.1388 0.0139 9.985
## ProcessorGPU:MatrixOperationInversion -106.4466 34.5267 -3.083
## ProcessorTPU:MatrixOperationInversion -130.3712 34.5267 -3.776
## ProcessorGPU:MatrixOperationMultiplication -89.5616 34.5267 -2.594
## ProcessorTPU:MatrixOperationMultiplication -90.9788 34.5267 -2.635
## Pr(>|t|)
## (Intercept) 0.019810 *
## ProcessorGPU 0.922179
## ProcessorTPU 0.919660
## MatrixOperationInversion 1.69e-07 ***
## MatrixOperationMultiplication 0.000227 ***
## MatrixSize < 2e-16 ***
## ProcessorGPU:MatrixOperationInversion 0.002212 **
## ProcessorTPU:MatrixOperationInversion 0.000187 ***
## ProcessorGPU:MatrixOperationMultiplication 0.009886 **
## ProcessorTPU:MatrixOperationMultiplication 0.008787 **
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 109.2 on 350 degrees of freedom
## Multiple R-squared: 0.3218, Adjusted R-squared: 0.3043
## F-statistic: 18.45 on 9 and 350 DF, p-value: < 2.2e-16
```

```
#res <- resid(mod2NI)
#plot(fitted(mod2NI), res)
#abline(0,0)

#boxcox transform
boxCox(mod2N, family="yjPower", plotit = TRUE)
```



```
lambda=-1
depvar.transformed <- yjPower(data$Runtime, lambda)
Processor <- data$Processor
MatrixOperation <- data$MatrixOperation
MatrixSize <- data$MatrixSize
mod2NT<-lm(depvar.transformed~Processor+MatrixOperation+MatrixSize)
summary(mod2NT)
```

```
##
## Call:
## lm(formula = depvar.transformed ~ Processor + MatrixOperation +
##      MatrixSize)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.56272 -0.17233 -0.05672  0.21517  0.50449
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.859e-01  2.873e-02   6.472 3.22e-10 ***
## ProcessorGPU   -6.921e-02  2.975e-02  -2.326 0.020587 *
## ProcessorTPU   -4.210e-01  2.975e-02 -14.148 < 2e-16 ***
## MatrixOperationInversion  3.280e-01  2.975e-02  11.023 < 2e-16 ***
## MatrixOperationMultiplication 1.107e-01  2.975e-02   3.719 0.000233 ***
## MatrixSize      3.815e-04  2.934e-05  13.006 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2305 on 354 degrees of freedom
## Multiple R-squared:  0.5974, Adjusted R-squared:  0.5917
## F-statistic: 105 on 5 and 354 DF, p-value: < 2.2e-16
```

```
#res <- resid(mod2NT)
#plot(fitted(mod2NT), res)
#abline(0,0)
```


Zhanhao Zhang Part II

General Visualization


Lixian Chen Part II: Plots

Interaction Plots


```
jpeg(filename = "../figs/interaction_size_time.jpeg", width = 600, height = 400, quality = 10000)
interaction.plot(df$MatrixSize, df$MatrixOperation, df$Runtime)
while (!is.null(dev.list())) dev.off()
```




```
jpeg(filename = "../figs/interaction_size_log_time.jpeg", width = 600, height = 400, quality = 10000)
interaction.plot(df$MatrixSize, df$MatrixOperation, log(df$Runtime))
while (!is.null(dev.list())) dev.off()
```




```
jpeg(filename = "../figs/interaction_CPU_size_time.jpeg", width = 600, height = 400, quality = 10000)
interaction.plot(CPUdata$MatrixSize, CPUdata$MatrixOperation, CPUdata$Runtime)
while (!is.null(dev.list())) dev.off()
```



```
jpeg(filename = "../figs/interaction_GPU_size_time.jpeg", width = 600, height = 400, quality = 10000)
interaction.plot(GPUdata$MatrixSize, GPUdata$MatrixOperation, GPUdata$Runtime)
while (!is.null(dev.list())) dev.off()
```




```
jpeg(filename = "../figs/interaction_TPU_size_time.jpeg", width = 600, height = 400, quality = 10000)
interaction.plot(TPUdata$MatrixSize, TPUdata$MatrixOperation, TPUdata$Runtime)
while (!is.null(dev.list())) dev.off()
```





Boxplots


```
jpeg(filename = "../figs/Operation_vs_runtime_size320.jpeg", width = 1400, height = 400, quality = 10000)
boxplot(Runtime~Processor*MatrixOperation, data = size320data, main="At the level of matrix size=320")
while (!is.null(dev.list())) dev.off()
```



```
jpeg(filename = "../figs/Operation_vs_runtime_size640.jpeg", width = 1400, height = 400, quality = 10000)
boxplot(Runtime~Processor*MatrixOperation, data = size640data, main="At the level of matrix size=640")
while (!is.null(dev.list())) dev.off()
```



```
jpeg(filename = "../figs/Operation_vs_runtime_size1280.jpeg", width = 1400, height = 400, quality = 10000)
boxplot(Runtime~Processor*MatrixOperation, data = size1280data, main="At the level of matrix size=1280")
while (!is.null(dev.list())) dev.off()
```




Jingbin Cao Part II

Plot Anova Residuals


CPU Add

```
jpeg(filename = "../figs/res_cpu_add.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_cpu_add$residuals))
plot(mod_cpu_add$fitted.values, mod_cpu_add$residuals, ylim=c(-m, m), main = "Residual Plot for CPU and GPU")
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_cpu_add.jpeg")
```




CPU Multiplication

```
jpeg(filename = "../figs/res_cpu_mult.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_cpu_mult$residuals))
plot(mod_cpu_mult$fitted.values, mod_cpu_mult$residuals, ylim=c(-m, m), main = "Residual Plot for CPU and Addition")
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_cpu_mult.jpeg")
```




CPU Inversion

```
jpeg(filename = "../figs/res_cpu_inv.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_cpu_inv$residuals))
plot(mod_cpu_inv$fitted.values, mod_cpu_inv$residuals, ylim=c(-m, m), main = "Residual Plot for CPU and Multiplication")
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_cpu_inv.jpeg")
```




GPU Add

```
jpeg(filename = "../figs/res_gpu_add.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_gpu_add$residuals))
plot(mod_gpu_add$fitted.values, mod_gpu_add$residuals, ylim=c(-m, m), main = "Residual Plot for GPU and Inversion")
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_gpu_add.jpeg")
```




GPU Multiplication

```
jpeg(filename = "../figs/res_gpu_mult.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_gpu_mult$residuals))
plot(mod_gpu_mult$fitted.values, mod_gpu_mult$residuals, ylim=c(-m, m), main = "Residual Plot for GPU and Addition")
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_gpu_mult.jpeg")
```




GPU Inversion

```
jpeg(filename = "../figs/res_gpu_inv.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_gpu_inv$residuals))
plot(mod_gpu_inv$fitted.values, mod_gpu_inv$residuals, ylim=c(-m, m), main = "Residual Plot for GPU and
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_gpu_inv.jpeg")
```




TPU Add

```
jpeg(filename = "../figs/res_tpu_add.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_tpu_add$residuals))
plot(mod_tpu_add$fitted.values, mod_tpu_add$residuals, ylim=c(-m, m), main = "Residual Plot for TPU and
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_tpu_add.jpeg")
```

TPU Multiplication


```
jpeg(filename = "../figs/res_tpu_mult.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_tpu_mult$residuals))
plot(mod_tpu_mult$fitted.values, mod_tpu_mult$residuals, ylim=c(-m, m), main = "Residual Plot for TPU and Addition")
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_tpu_mult.jpeg")
```



TPU Inversion

```
jpeg(filename = "../figs/res_tpu_inv.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(mod_tpu_inv$residuals))
plot(mod_tpu_inv$fitted.values, mod_tpu_inv$residuals, ylim=c(-m, m), main = "Residual Plot for TPU and
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_tpu_inv.jpeg")
```

Residual Plot for TPU and Inversion



Linear Regression for Matrix Size as continuous We just need to run linear regression on seven pairs that do have Matrix Size Effects

```
summary(lm_cpu_add <- lm(log10(Runtime)~as.numeric(MatrixSize),data=data[data$Processor == "CPU" & data$MatrixOperation == "Inversion",])
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ as.numeric(MatrixSize), data = data[data$Processor ==
## "CPU" & data$MatrixOperation == "Inversion", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.35107 -0.18405 -0.13000  0.08052  0.62584
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.5277230   0.0646279  -23.64  <2e-16 ***
## as.numeric(MatrixSize)  0.0024138  0.0001237   19.52  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.3239 on 38 degrees of freedom
## Multiple R-squared:  0.9093, Adjusted R-squared:  0.9069
## F-statistic: 380.9 on 1 and 38 DF,  p-value: < 2.2e-16

summary(lm_cpu_mult <- lm(log10(Runtime)~as.numeric(MatrixSize),data[data$Processor == "CPU" & data$Matr

##
## Call:
## lm(formula = log10(Runtime) ~ as.numeric(MatrixSize), data = data[data$Processor ==
##      "CPU" & data$MatrixOperation == "Multiplication", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6282 -0.5727 -0.1663  0.5620  0.9143
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -0.9187389   0.1195854   -7.683 2.99e-09 ***
## as.numeric(MatrixSize)  0.0033880   0.0002288  14.805 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5993 on 38 degrees of freedom
## Multiple R-squared:  0.8522, Adjusted R-squared:  0.8484
## F-statistic: 219.2 on 1 and 38 DF,  p-value: < 2.2e-16

summary(lm_cpu_inv <- lm(log10(Runtime)~as.numeric(MatrixSize),data[data$Processor == "CPU" & data$Matr

##
## Call:
## lm(formula = log10(Runtime) ~ as.numeric(MatrixSize), data = data[data$Processor ==
##      "CPU" & data$MatrixOperation == "Inversion", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.79157 -0.41955 -0.03787  0.47617  0.67656
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -0.324820   0.100347   -3.237 0.00251 **
## as.numeric(MatrixSize)  0.002947   0.000192  15.346 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5029 on 38 degrees of freedom
## Multiple R-squared:  0.8611, Adjusted R-squared:  0.8574
## F-statistic: 235.5 on 1 and 38 DF,  p-value: < 2.2e-16

summary(lm_gpu_add <- lm(log10(Runtime)~as.numeric(MatrixSize),data=data[data$Processor == "GPU" & data$Matr

##
```

```
## Call:
## lm(formula = log10(Runtime) ~ as.numeric(MatrixSize), data = data[data$Processor ==
##      "GPU" & data$MatrixOperation == "Addition", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.22827 -0.05771  0.04545  0.06655  0.14214
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.2307302   0.0231517  -53.16  <2e-16 ***
## as.numeric(MatrixSize)  0.0006241  0.0000443   14.09  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.116 on 38 degrees of freedom
## Multiple R-squared:  0.8393, Adjusted R-squared:  0.835
## F-statistic: 198.4 on 1 and 38 DF,  p-value: < 2.2e-16
```

```
summary(lm_gpu_mult <- lm(log10(Runtime)~as.numeric(MatrixSize),data[data$Processor == "GPU" & data$Matr
```

```
##
## Call:
## lm(formula = log10(Runtime) ~ as.numeric(MatrixSize), data = data[data$Processor ==
##      "GPU" & data$MatrixOperation == "Multiplication", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.189553 -0.042371 -0.006256  0.062402  0.136090
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -1.148e+00  1.822e-02  -63.01  <2e-16 ***
## as.numeric(MatrixSize)  1.715e-03  3.487e-05   49.19  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09131 on 38 degrees of freedom
## Multiple R-squared:  0.9845, Adjusted R-squared:  0.9841
## F-statistic: 2419 on 1 and 38 DF,  p-value: < 2.2e-16
```

```
summary(lm_gpu_inv <- lm(log10(Runtime)~as.numeric(MatrixSize),data[data$Processor == "GPU" & data$Matr
```


```
##
## Call:
## lm(formula = log10(Runtime) ~ as.numeric(MatrixSize), data = data[data$Processor ==
##      "GPU" & data$MatrixOperation == "Inversion", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.36329 -0.27442  0.00045  0.24043  0.39749
##
## Coefficients:
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.3755794  0.0557293   6.739 5.55e-08 ***
## as.numeric(MatrixSize) 0.0015252  0.0001066  14.301 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2793 on 38 degrees of freedom
## Multiple R-squared:  0.8433, Adjusted R-squared:  0.8392
## F-statistic: 204.5 on 1 and 38 DF,  p-value: < 2.2e-16
```

Plot LM


CPU Addition

```
jpeg(filename = "../figs/res_lm_cpu_add.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(lm_cpu_add$residuals))
plot(lm_cpu_add$fitted.values, lm_cpu_add$residuals, ylim=c(-m, m), main = "Residual Plot for CPU and A
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_lm_cpu_add.jpeg")
```




CPU Multiplication

```
jpeg(filename = "../figs/res_lm_cpu_mult.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(lm_cpu_mult$residuals))
plot(lm_cpu_mult$fitted.values, lm_cpu_mult$residuals, ylim=c(-m, m), main = "Residual Plot for CPU and
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_lm_cpu_mult.jpeg")
```




CPU Inversion

```
jpeg(filename = "../figs/res_lm_cpu_inv.jpeg", width = 600, height = 400,quality = 10000)
m <- max(abs(lm_cpu_inv$residuals))
plot(lm_cpu_inv$fitted.values, lm_cpu_inv$residuals, ylim=c(-m, m), main = "Residual Plot for CPU and I
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_lm_cpu_inv.jpeg")
```


GPU Addition

```
jpeg(filename = "../figs/res_lm_gpu_add.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(lm_gpu_add$residuals))
plot(lm_gpu_add$fitted.values, lm_gpu_add$residuals, ylim=c(-m, m), main = "Residual Plot for GPU and A",
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_lm_gpu_add.jpeg")
```




GPU Multiplication

```
jpeg(filename = "../figs/res_lm_gpu_mult.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(lm_gpu_mult$residuals))
plot(lm_gpu_mult$fitted.values, lm_gpu_mult$residuals, ylim=c(-m, m), main = "Residual Plot for GPU and
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_lm_gpu_mult.jpeg")
```




GPU Inversion

```
jpeg(filename = "../figs/res_lm_gpu_inv.jpeg", width = 600, height = 400, quality = 10000)
m <- max(abs(lm_gpu_inv$residuals))
plot(lm_gpu_inv$fitted.values, lm_gpu_inv$residuals, ylim=c(-m, m), main = "Residual Plot for GPU and I
abline(h=0, col="red")
while (!is.null(dev.list())) dev.off()
knitr::include_graphics("../figs/res_lm_gpu_inv.jpeg")
```




At the level of matrix size=1280, avoid using CPU for inversion and multiplication because its run-times are much bigger.


For CPU and Addition with Different Matrix Size




For CPU and Multiplication with Different Matrix Size




For CPU and Inversion with Different Matrix Size




For GPU and Addition with Different Matrix Size




For GPU and Multiplication with Different Matrix Size




For GPU and Inversion with Different Matrix Size



For TPU and Addition with Different Matrix Size



For TPU and Multiplication with Different Matrix Size



For TPU and Inverction with Different Matrix Size

